

- Advanced AI-Driven Context-Aware Anonymization for Social Media Privacy Protection: A Survey
 - Abstract
 - I. INTRODUCTION
 - A. Background and Motivation
 - B. Privacy Protection Evolution
 - C. Motivation for Context-Aware Approaches
 - D. Survey Scope
 - II. LITERATURE SURVEY
 - A. Traditional Privacy Protection Approaches
 - A.1 Pixelation and Blurring
 - A.2 Masking Techniques
 - A.3 Face Detection-Based Anonymization
 - B. Deep Learning for Object Detection
 - B.1 YOLO (You Only Look Once) Architecture
 - B.2 Faster R-CNN and Region-Based Approaches
 - B.3 Single-Shot MultiBox Detectors (SSD)
 - B.4 EfficientNet and Lightweight Models
 - C. Specialized Detection Models
 - C.1 Face Detection Models
 - C.2 License Plate Detection and Recognition
 - C.3 Document Detection and Recognition
 - D. Optical Character Recognition (OCR)
 - D.1 Text Detection
 - D.2 Text Recognition
 - D.3 Named Entity Recognition (NER) for Sensitive Text
 - E. Context-Aware Image Processing
 - E.1 Semantic Segmentation
 - E.2 Saliency Detection
 - E.3 Foreground-Background Separation
 - F. Privacy Risk Assessment Frameworks
 - F.1 Privacy Scoring Models
 - F.2 Information-Theoretic Privacy
 - F.3 User-Centric Privacy Models
 - G. Privacy-Preserving Architectures
 - G.1 Edge Computing and On-Device Processing
 - G.2 Model Compression Techniques

- G.3 Federated Learning
- H. Challenges and Limitations in Current Approaches
 - H.1 Detection Challenges
 - H.2 Video Processing Challenges
 - H.3 User Control and Customization
 - H.4 Quality-Privacy Trade-offs
 - H.5 Interpretability and Explainability
- I. Emerging Trends and Future Directions
 - I.1 Vision Transformers and Multi-Modal Approaches
 - I.2 Adversarial Robustness
 - I.3 Regulatory Compliance
 - I.4 Generative Models for Realistic Replacement
- II-A Project Literature Summary (Team Review)
 - Background
 - Problem Statement
 - Research Gap
 - Proposed Solution (summary)
 - Selected Papers (team review highlights)
 - Research Gaps Identified by Team
 - Finalized Project Objectives (team)
 - Proposed Methodology (team summary)
 - Team References
- III. SYSTEM DESIGN AND ARCHITECTURAL CONSIDERATIONS
 - A. Conceptual System Architecture
 - B. Key Architectural Considerations from Literature
 - B.1 Privacy-First Design Principle
 - B.2 Computational Efficiency
 - B.3 Quality vs. Privacy Balance
 - B.4 Robustness and Reliability
 - B.5 Scalability and Extensibility
 - C. Anonymization Techniques Discussed in Literature
 - C.1 Traditional Techniques
 - C.2 Advanced Quality-Preserving Techniques
 - C.3 Adaptive and Content-Aware Approaches
 - D. Privacy Risk Assessment Approaches
 - D.1 Multi-Factor Risk Assessment
 - D.2 Information-Theoretic Frameworks
 - D.3 User-Centric Risk Models

- E. Implementation Considerations from Literature
 - E.1 Detection Model Integration
 - E.2 Performance Optimization Strategies
 - E.3 User Interface Design
- IV. COMPARATIVE ANALYSIS AND FINDINGS
 - A. Comparison of Privacy Protection Approaches
 - B. Context-Aware vs. Blanket Anonymization
 - C. Privacy-Quality Trade-offs
- V. CHALLENGES, OPEN PROBLEMS, AND FUTURE DIRECTIONS
 - A. Current Research Challenges
 - A.1 Technical Challenges
 - A.2 User Experience Challenges
 - A.3 Privacy and Security Challenges
 - B. Emerging Research Directions
 - B.1 Advanced Deep Learning Architectures
 - B.2 Privacy-Enhancing Technologies
 - B.3 Explainable and Trustworthy Privacy Systems
 - B.4 Generative Models for Privacy
 - B.5 Regulatory Compliance and Privacy Standards
 - B.6 Behavioral and Contextual Privacy
- VI. CONCLUSIONS AND FUTURE WORK
 - A. Key Findings
 - B. Open Research Questions
 - C. Recommended Research Directions
 - D. Societal Implications
 - E. Final Remarks
- REFERENCES

Advanced AI-Driven Context-Aware Anonymization for Social Media Privacy Protection: A Survey

Prof Nidhishree G

Assistant Professor

Department of CSE(AI&ML)

Vidyavardhaka College of Engineering

Team Members

- DARSHAN I C (4VV23CI014) - 4vv23ci014@vvce.ac.in
- MOULYA SHREE V (4VV23CI053) - 4vv23ci053@vvce.ac.in
- MADHUSUDAN S (4VV24CI404) - 4vv24ci404@vvce.ac.in

Department of CSE(AI&ML)
Vidyavardhaka College of Engineering
Mysuru, India

Abstract

Social media platforms have revolutionized digital communication, enabling billions of users to share content globally. However, this widespread sharing has created unprecedented privacy challenges, with users frequently exposing sensitive information including faces, license plates, identity documents, and personal text. This comprehensive survey reviews state-of-the-art techniques for protecting personal information in social media content through AI-driven anonymization. The paper examines existing privacy protection methodologies, ranging from traditional image manipulation techniques (pixelation, blurring, masking) to advanced deep learning-based approaches. We analyze contemporary object detection models (YOLO, Faster R-CNN, RetinaFace), facial recognition technologies, Optical Character Recognition (OCR), and semantic understanding methods. This survey compares context-aware selective anonymization approaches with conventional blanket-blurring techniques, highlighting the trade-off between privacy preservation and image quality maintenance. We review privacy risk assessment frameworks, user preference modeling, and emerging trends in local, privacy-first architectures. Through systematic analysis of 50+ relevant research publications, this paper identifies current limitations in existing systems and discusses future research directions including federated learning, adversarial robustness, and regulatory compliance mechanisms. The survey provides a foundation for understanding next-generation privacy protection systems that balance user control, computational efficiency, and strong privacy guarantees.

Keywords: Privacy Protection, AI Anonymization, Deep Learning, Object Detection, Face Detection, OCR, Context-Aware Processing, Privacy Risk Assessment, Social Media Privacy, Image Processing

I. INTRODUCTION

A. Background and Motivation

The exponential growth of social media platforms has fundamentally transformed how individuals share information and interact globally. According to recent statistics, over 4.9 billion social media users worldwide actively share photos, videos, and personal updates daily across platforms such as Facebook, Instagram, Twitter, TikTok, and emerging platforms [1][2]. This unprecedented volume of shared content has created significant privacy challenges that extend beyond individual users to impact organizations, governments, and society at large.

Research by Pew Research Center and Statista indicates that approximately 64% of social media users have inadvertently exposed personal information through shared content [3]. Common examples of sensitive information leakage include:

- **Facial Information:** Profile pictures, family photos, and event images revealing identity and recognizable features
- **Location Data:** Geotags, landmark recognition, and environmental context enabling physical location tracking
- **Identity Documents:** Accidental exposure of driver's licenses, passports, and government-issued IDs in images
- **Personal Text:** Phone numbers, addresses, email addresses, and names visible in images
- **License Plates:** Vehicle identification information visible in street or parking images
- **Biometric Data:** Facial geometry, gait patterns, and voice characteristics from multimedia content

The consequences of such privacy breaches are far-reaching, including identity theft, location-based harassment, social engineering attacks, unauthorized surveillance, and data aggregation for discriminatory purposes [4][5].

B. Privacy Protection Evolution

Traditional approaches to privacy protection have evolved significantly over the past two decades:

1. **Manual Redaction Era (2000s):** Users manually edited photos using basic image editing tools, a labor-intensive and error-prone process
2. **Automated Blurring Era (2010s):** Platforms introduced automated face blurring and basic anonymization features, though often with poor quality and limited object types
3. **Context-Unaware Era (2015-2020):** Early deep learning systems provided face detection but lacked contextual understanding, applying uniform anonymization
4. **Context-Aware Era (2020-present):** Advanced systems attempt to understand scene context and selectively protect information based on content analysis

C. Motivation for Context-Aware Approaches

Conventional blanket-blurring approaches, while effective at removing sensitive information, introduce significant drawbacks that motivate research into more sophisticated solutions:

- **Visual Quality Degradation:** Indiscriminate blurring reduces image utility and aesthetic value [6]
- **Information Loss:** Necessary contextual information is removed along with sensitive elements
- **User Dissatisfaction:** Over-aggressive anonymization discourages content sharing [7]
- **Computational Inefficiency:** Processing entire images when only small regions require protection

Contemporary research emphasizes the need for intelligent, context-aware approaches that:

- Selectively protect only genuinely sensitive information
- Preserve image quality and contextual richness
- Provide user control and customization
- Offer transparency and explainability in privacy decisions
- Implement privacy-first, local processing architectures

D. Survey Scope

This survey comprehensively reviews the state-of-the-art in AI-driven privacy protection for social media content. Our analysis covers:

1. Traditional Privacy Protection Techniques
2. Deep Learning for Object Detection
3. Specialized Detection Models (Face, Text, Documents)
4. Context-Aware Processing Methods
5. Privacy Risk Assessment Frameworks
6. Privacy-Preserving Architectures
7. Challenges and Open Problems

The paper is organized as follows: Section II presents an extensive literature survey of relevant techniques and methodologies. Section III discusses system design concepts and architectural considerations from surveyed literature. Section IV summarizes findings from relevant research, presenting comparative analyses. Section V identifies research gaps, future directions, and emerging challenges.

II. LITERATURE SURVEY

A. Traditional Privacy Protection Approaches

Conventional image anonymization techniques have been extensively studied in computer vision and privacy literature. Early approaches employed simple image manipulation techniques:

A.1 Pixelation and Blurring

The most widely used traditional method applies uniform pixelation or Gaussian blur across predefined regions [8][9]. While simple to implement, these methods lack sophistication and indiscriminately reduce image quality. The approach divides images into blocks of varying sizes and replaces pixel values with block averages for pixelation, or applies Gaussian kernels for smooth blurring.

Research by Gross et al. [10] demonstrated that even pixelated faces could be re-identified through advanced computational techniques, questioning the effectiveness of blanket approaches. Studies have shown that faces pixelated with blocks as small as 8×8 pixels could be reconstructed using super-resolution techniques [11].

Advantages:

- Simple implementation requiring minimal computational resources
- Universally applicable to any image content
- Reversible with known parameters (security risk)

Disadvantages:

- Crude approach reducing visual quality significantly
- Fails to provide formal privacy guarantees
- Vulnerability to reversal attacks and re-identification

A.2 Masking Techniques

Physical masking through rectangles, circles, or custom shapes provides more control than uniform blurring [12]. Masks can be solid colors, patterns, or even replacement images. These techniques remain visually distracting and provide no context awareness.

Variants include:

- **Rectangular Masking:** Simple bounding box replacement
- **Circular Masking:** Circular region covering for faces
- **Arbitrary Shape Masking:** Custom polygon-based masking

Research findings:

- Masks larger than necessary reduce image utility [13]
- User studies show masks are visually jarring [14]
- Incomplete masking leaves identifiable fragments [15]

A.3 Face Detection-Based Anonymization

Leveraging traditional face detection algorithms (Haar Cascades, Active Shape Models), researchers developed automated face-blurring systems [16]. The limitation of

these approaches is their inability to detect other sensitive information types, such as license plates, documents, or personal text.

Classic algorithms:

- **Haar Cascade Classifiers:** Fast but prone to false positives [17]
- **Active Shape Models:** Better accuracy but computationally expensive [18]
- **Deformable Part Models:** Improved robustness [19]

Shortcomings:

- Limited to face detection only
- High false positive rates (10-30% depending on conditions) [20]
- Poor performance with occlusion, extreme poses, or poor lighting
- No privacy risk quantification

B. Deep Learning for Object Detection

Recent advances in deep learning have significantly improved object detection accuracy and speed, enabling real-time privacy protection applications:

B.1 YOLO (You Only Look Once) Architecture

Introduced by Redmon et al. [21], YOLO represents a paradigm shift in object detection by treating detection as a single regression problem. Rather than generating region proposals, YOLO processes entire images in one pass, achieving detection speeds of 45-155 fps depending on model variants (YOLOv3, YOLOv4, YOLOv5, YOLOv8).

Key characteristics:

- **Real-time Performance:** Processes full images in single forward pass [22]
- **Architecture Evolution:**
 - YOLOv1-v3: Original single-shot detectors [21][23]
 - YOLOv4: Enhanced backbone and feature extraction [24]
 - YOLOv5: Improved accuracy with reduced model size [25]
 - YOLOv8: Latest architecture with edge optimization [26]

Performance metrics from literature:

- License plate detection: 92-94% accuracy [27]
- General object detection: COCO mAP @0.5 up to 56% [28]

- Processing speed: 25-30 fps on CPU, 100+ fps on GPU [29]

Studies comparing to alternatives:

- Superior performance vs R-CNN variants when applied to license plate detection [30]
- Better speed-accuracy trade-off for real-time applications [31]

B.2 Faster R-CNN and Region-Based Approaches

These region-based convolutional neural networks [32] achieve high accuracy through a two-stage detection process: region proposal generation followed by classification and refinement. The architecture processes regions of interest (RoIs) through fully connected layers.

Variants and extensions:

- **Faster R-CNN:** Original two-stage detector [32]
- **Mask R-CNN:** Adds instance segmentation [33]
- **Cascade R-CNN:** Improved IoU matching [34]

Characteristics:

- While more accurate than YOLO in certain scenarios, requiring significantly more computational resources [35]
- Limiting real-time applicability on mobile devices [36]
- Better for detecting small objects [37]
- Average inference time: 100-300ms per image on CPU [38]

B.3 Single-Shot MultiBox Detectors (SSD)

SSD architectures [39] provide a compromise between YOLO's speed and Faster R-CNN's accuracy by using multi-scale feature maps.

Key contributions:

- Multi-scale feature maps for detection at different scales [39]
- Convolutional predictors for object detection [40]
- Competitive performance with lower computational cost [41]

B.4 EfficientNet and Lightweight Models

Lightweight architectures [42] enable edge device deployment, crucial for privacy-first local processing. EfficientNet particularly provides excellent accuracy-efficiency trade-offs through compound model scaling.

Family of models:

- **EfficientNet-B0 to B7:** Scaling up parameters systematically [42]
- **MobileNet:** Depthwise separable convolutions for efficiency [43]
- **SqueezeNet:** Extreme model compression [44]
- **ShuffleNet:** Channel shuffling for mobile [45]

Deployment considerations:

- EfficientNet-B0: 5.3M parameters, 390M FLOPs
- MobileNetV2: 3.5M parameters, 300M FLOPs
- Achievable 30+ fps inference on mobile devices [46]

C. Specialized Detection Models

C.1 Face Detection Models

Specialized architectures achieve state-of-the-art accuracy in face detection across diverse conditions including partial occlusion, varying illumination, and extreme poses:

Modern Face Detection Approaches:

1. RetinaFace Architecture [47]

- Multi-task learning combining face detection, landmark detection, and 3D face reconstruction
- Accuracy: >99% on benchmark datasets [47]
- Handles challenging scenarios: extreme poses, illumination, occlusion
- Processing speed: 15-20 fps on CPU

2. MTCNN (Multi-task Cascaded CNNs) [48]

- Three-stage cascaded architecture for face detection and alignment
- Accuracy: 96-98% on benchmark datasets [48]
- Robust to significant variations in pose and scale
- Computational cost: slower than modern approaches

3. MediaPipe Face Detection [49]

- Lightweight, mobile-optimized face detection
- Performance: 30+ fps on mobile devices [49]
- Designed for edge deployment with accuracy trade-offs
- Suitable for real-time mobile applications

4. S3FD (Single Shot Scale-Invariant Face Detector) [50]

- Designed for detecting small faces at various scales
- Accuracy: 96%+ on challenging datasets
- Effective for wide-view images with multiple faces

Comparative analysis:

- RetinaFace: Best accuracy but higher computational cost
- MTCNN: Robust but slower inference
- MediaPipe: Optimal for mobile deployment
- S3FD: Superior small face detection [51]

C.2 License Plate Detection and Recognition

License plate detection is a specialized computer vision problem requiring robust object detection followed by text recognition:

Detection approaches:

- **YOLO-based Methods:** Trained on license plate datasets, achieving 92-96% accuracy [52]
- **SSD variants:** Modified for plate detection with competitive performance [53]
- **Custom architectures:** Domain-specific optimizations [54]

Recognition techniques:

- **Traditional OCR:** Tesseract, commercial systems achieving 85-90% accuracy [55]
- **Deep Learning OCR:** CRNN achieving 91-95% accuracy [56]
- **Sequence models:** Attention-based recognition for variable-length plates [57]

Performance metrics from literature:

- Detection accuracy: 92-96% under normal conditions [52]

- Recognition accuracy: 90-95% for clear plates [56]
- Degradation under challenging conditions: 20-40% accuracy drop [58]

C.3 Document Detection and Recognition

Identity documents require specialized detection due to their distinct visual properties:

Document detection approaches:

- **YOLO-based training:** Custom datasets for document bounding box detection
- **Perspective correction:** Handling document tilt and perspective distortion [59]
- **Layout analysis:** Understanding document structure and regions [60]

Document types commonly detected:

- Passports and travel documents
- Driver's licenses
- ID cards
- Birth certificates
- Financial documents

Recognition techniques:

- **Field extraction:** Detecting specific regions (photo, signature, text) [61]
- **OCR on documents:** Extracting text with field context [62]
- **Microprint detection:** Identifying security features [63]

D. Optical Character Recognition (OCR)

OCR technology enables detection and protection of sensitive textual information. Recent deep learning-based OCR systems employ advanced techniques:

D.1 Text Detection

CRAFT (Character Region Awareness For Text Detection) and EAST (Efficient and Accurate Scene Text Detector) are state-of-the-art approaches:

CRAFT Architecture [64]

- Character-level region detection providing character bounding boxes
- Accuracy: >95% on standard benchmarks [64]

- Robust to various text orientations and scales
- Processing speed: 10-15 fps on CPU

EAST Architecture [65]

- Single neural network predicting text regions directly
- Simplicity and efficiency compared to prior methods
- Accuracy: 94-96% depending on dataset [65]
- Processing speed: Similar to CRAFT

Other approaches:

- **TextSnake:** Capturing arbitrary-shaped text [66]
- **PixelLink:** Pixel-level text region connectivity [67]
- **PSENet:** Progressive scale expansion network [68]

D.2 Text Recognition

CRNN (Convolutional Recurrent Neural Network) architectures achieve 90%+ accuracy:

CRNN Structure [69]

- Combines CNN for feature extraction with RNN for sequence modeling
- Handles variable-length text sequences [69]
- Accuracy: 90-93% on standard benchmarks [70]

Recent improvements:

- **Attention mechanisms:** Context-aware character prediction [71]
- **Transformer-based approaches:** Superior sequence understanding [72]
- **End-to-end learning:** Joint detection and recognition [73]

D.3 Named Entity Recognition (NER) for Sensitive Text

Advanced NLP models identify sensitive text categories within recognized text:

Categories typically protected:

- Personal names
- Addresses and locations
- Phone numbers

- Email addresses
- Financial information
- Social security numbers
- Medical information

NER approaches:

- **BiLSTM-CRF models:** Widely used for NER tasks [74]
- **BERT-based models:** Superior context understanding [75]
- **Domain-specific models:** Trained on sensitive data patterns [76]

Performance metrics:

- Standard NER F1-scores: 85-95% depending on domain [77]
- Sensitive information detection: 88-94% accuracy [78]

E. Context-Aware Image Processing

Understanding scene context represents a frontier in intelligent image anonymization:

E.1 Semantic Segmentation

FCN (Fully Convolutional Networks) and U-Net architectures enable pixel-level classification:

Key architectures:

- **FCN (Fully Convolutional Networks)** [79]
 - First end-to-end semantic segmentation network
 - Baseline accuracy: 62-65% on Pascal VOC [79]
 - Pioneering work enabling pixel-level predictions
- **U-Net** [80]
 - Encoder-decoder structure with skip connections
 - Accuracy: 65-75% on standard benchmarks [80]
 - Originally designed for biomedical image segmentation
- **DeepLab Series** [81][82][83]
 - Atrous convolutions for multi-scale context

- DeepLabv3+: 78-79% mIoU on Cityscapes [83]
- Widely adopted for scene understanding
- **Transformer-based approaches:**
 - **Segformer:** Vision transformer for segmentation [84]
 - **SETR:** Semantic segmentation with transformers [85]
 - Superior context understanding vs CNNs [86]

Utility for privacy:

- Distinguishing foreground subjects from background [87]
- Identifying document boundaries and regions [88]
- Scene understanding for context-aware anonymization [89]

E.2 Saliency Detection

Computational models of visual attention identify image regions that attract human visual focus:

Saliency models:

- **Traditional approaches:** Contrast, center-bias, and feature integration [90]
- **Deep learning saliency:** CNN-based attention prediction [91]
- **Eye-tracking inspired:** Models trained on gaze patterns [92]

Application to privacy:

- Prioritizing anonymization of visible sensitive objects [93]
- De-prioritizing less visible sensitive objects
- Balancing privacy and quality based on visual importance [94]

E.3 Foreground-Background Separation

Advanced techniques for separating scene elements:

Approaches:

- **Matting techniques:** Alpha blending for smooth boundaries [95]
- **Segmentation-based:** Using semantic maps for separation [96]
- **Learning-based:** Deep networks for robust separation [97]

Privacy implications:

- Preserving main subject while protecting background privacy [98]
- Selective anonymization based on object importance

F. Privacy Risk Assessment Frameworks

Quantifying privacy risk remains a challenging research problem:

F.1 Privacy Scoring Models

Shmueli et al. [99] developed metrics combining multiple factors:

Multi-factor approaches:

- Number of sensitive objects detected
- Type of sensitive information (faces > documents > plates)
- Visibility and prominence of objects
- Metadata presence (location, timestamp)
- Uniqueness and re-identification risk [100]

Challenges:

- Lack of standardization across systems [101]
- Context dependency of privacy definitions [102]
- Individual variation in privacy preferences [103]

F.2 Information-Theoretic Privacy

Formal privacy quantification through theoretical frameworks:

Differential Privacy [104]

- Formal privacy guarantee with epsilon parameter
- Quantifies maximum privacy loss [105]
- Implementation challenges: accuracy-privacy trade-offs [106]

Information-Theoretic Approaches [107]

- Mutual information reduction metrics
- Shannon entropy calculations

- Formal but computationally intensive [108]

F.3 User-Centric Privacy Models

Research demonstrates significant variation in individual privacy preferences [109]:

Contextual Privacy [110]

- Privacy concerns vary by context and audience
- Social media posts have different privacy requirements than medical images
- Context-aware privacy definitions [111]

Privacy Preference Elicitation [112]

- Methods for understanding user privacy requirements
- Systems incorporating preferences achieve better user satisfaction [113]

G. Privacy-Preserving Architectures

Growing concern regarding data collection has motivated development of on-device processing approaches:

G.1 Edge Computing and On-Device Processing

Deploying deep learning models on edge devices (smartphones, IoT devices) eliminates cloud transmission requirements [114]:

Benefits:

- Data never leaves user's device
- Immediate processing without network latency
- Works in offline scenarios
- Regulatory compliance (GDPR, CCPA) [115]

Challenges:

- Limited computational resources
- Model size and optimization requirements [116]
- Accuracy preservation with resource constraints [117]

Enabling technologies:

- **TensorFlow Lite:** Mobile inference framework [118]
- **ONNX Runtime:** Cross-platform inference [119]
- **Core ML:** Apple's machine learning framework [120]

G.2 Model Compression Techniques

Enabling edge deployment while preserving accuracy:

Quantization [121]

- INT8 quantization reduces model size by 75% with <2% accuracy loss
- Post-training quantization vs. quantization-aware training
- Impact on model compression and speed [122]

Pruning [123]

- Removing redundant weights and neurons
- Structured vs. unstructured pruning
- Achieves 10-50x compression with minimal accuracy loss [124]

Knowledge Distillation [125]

- Training smaller models to mimic larger ones
- Student-teacher learning paradigms
- Effective for deployment on resource-constrained devices [126]

Network Architecture Search [127]

- Automatically designing efficient architectures
- Finding optimal speed-accuracy trade-offs
- Mobile-optimized model discovery [128]

G.3 Federated Learning

Distributed model training without centralizing user data [129]:

Advantages for privacy:

- User data never leaves device [130]
- Model updates aggregated centrally
- Enables continuous improvement without privacy sacrifice [131]

Research contributions:

- Convergence analysis of federated learning [132]
- Communication efficiency [133]
- Privacy guarantees with differential privacy [134]

Challenges:

- Statistical heterogeneity across devices [135]
- Communication overhead [136]
- System and statistical heterogeneity [137]

H. Challenges and Limitations in Current Approaches

Despite significant progress, existing systems face several limitations:

H.1 Detection Challenges

Complex Environments

- Performance degrades under extreme lighting [138]
- Occlusion and partial visibility reduce accuracy [139]
- Motion blur and image noise impact detection [140]

Small Object Detection

- Face detection for distant persons [141]
- License plates at various scales [142]
- Document detection in cluttered scenes [143]

Language and Script Variations

- OCR performance varies across languages [144]
- Non-Latin scripts pose challenges [145]
- Historical documents and unusual fonts [146]

H.2 Video Processing Challenges

- Frame-by-frame processing inefficiency [147]
- Temporal consistency across frames [148]

- Flicker and visual artifacts in anonymized videos [149]

H.3 User Control and Customization

- Limited existing systems for user preference specification [150]
- One-size-fits-all approaches unsatisfactory [151]
- Need for intuitive preference interfaces [152]

H.4 Quality-Privacy Trade-offs

- Fundamental trade-off between strong privacy and image quality [153]
- User preferences vary regarding acceptable quality loss [154]
- Adaptive approaches needed for different content types [155]

H.5 Interpretability and Explainability

- Deep learning models often lack transparency [156]
- Users unable to understand anonymization decisions [157]
- Need for explainable AI approaches in privacy systems [158]

I. Emerging Trends and Future Directions

I.1 Vision Transformers and Multi-Modal Approaches

- Vision transformers for improved scene understanding [159]
- Multi-modal models combining text, image, and metadata [160]
- Attention mechanisms for context-aware processing [161]

I.2 Adversarial Robustness

- Robustness to adversarial attacks and edge cases [162]
- Certified defenses for privacy protection [163]
- Evaluation frameworks for privacy system robustness [164]

I.3 Regulatory Compliance

- GDPR, CCPA, and emerging privacy regulations [165]
- Automated compliance checking [166]
- Privacy-by-design frameworks [167]

I.4 Generative Models for Realistic Replacement

- GAN-based realistic face replacement [168]
 - Diffusion models for content-aware inpainting [169]
 - Maintaining naturalness while removing sensitive information [170]
-

II-A Project Literature Summary (Team Review)

Background

Social media users unknowingly expose sensitive personal information — faces, license plates, identity documents, house numbers, and personal text — in publicly shared images and videos.

Problem Statement

- Traditional privacy tools apply blanket blurring to all detected objects, degrading image quality and removing important visual context that users want to preserve.
- Existing systems lack context-awareness: they blur regardless of whether the face is the main subject, a background person, or even an image on a poster.

Research Gap

- No current system combines context-aware selective anonymization with Privacy Risk Scoring, user-controlled rules, and local privacy-first processing in a single unified pipeline.

Proposed Solution (summary)

The Social Media Privacy Guard system proposes using deep learning (YOLOv8, CNN, OCR) to detect sensitive objects and intelligently anonymize only relevant

background elements, preserving the primary subject while generating a measurable Privacy Risk Score (0–100).

Selected Papers (team review highlights)

The team's review tabulated key works, methods, advantages and limitations.

Highlights include:

- Frome et al. (2009) — Google Street View face and plate detection; high recall but domain-specific and day-time bias.
- Redmon et al. (2016) — YOLO: real-time unified detection; fast but lower precision on small/grouped objects.
- Zhang et al. (2016) — MTCNN: cascaded face detection and alignment; accurate but degrades on extreme occlusions.
- Hukkelås & Lindseth (2023) — DeepPrivacy2: GAN-based full-body anonymization; realistic but computationally expensive.
- Baek et al. (2019) — CRAFT: robust scene text detection; excels on arbitrary-shaped text but slower and may miss very small/low-contrast text.
- Li et al. (2018) — End-to-end license plate detection & recognition; unified pipeline with performance drops on occluded/distorted plates.
- Ren et al. (2016) — Faster R-CNN: region-proposal detectors with high accuracy; slower than single-shot detectors.
- Viola & Jones (2001) — Haar cascades: historically important, fast, high false positives in unconstrained settings.

Research Gaps Identified by Team

- Blanket blurring without context awareness leading to loss of main-subject quality.
- No main-subject preservation in existing tools.
- Reversible anonymization vulnerability: blur/pixelation can be reversed by modern AI.
- Limited OCR robustness for stylized/scene text.
- Lack of measurable Privacy Risk quantification in user-facing tools.

Finalized Project Objectives (team)

1. Intelligent Sensitive Object Detection — faces, plates, documents, house numbers, scene text (YOLOv8, CNNs, OCR).
2. Context-Aware Selective Anonymization — preserve main subject, anonymize background sensitive elements.
3. Privacy Risk Scoring — develop a 0–100 exposure metric.
4. Robust OCR-Based Text Privacy — integrate CRAFT + EasyOCR for stylized text.
5. User-Controlled Privacy Rules — toggle categories and blur strength.
6. Privacy-First Local Processing — on-device inference, no cloud upload.

Proposed Methodology (team summary)

Stage 1: Input processing — accept user image/video, preprocess (resize, normalize).

Stage 2: Multi-model detection — YOLOv8 (faces, plates, documents), CRAFT + EasyOCR (scene text), MTCNN/DeepFace as secondary face detector.

Stage 3: Context analysis — main-subject check (e.g., face bounding box >30% of frame), confidence filtering.

Stage 4: Anonymization — user-selected modes: Gaussian blur, pixelation, emoji overlay, or generative inpainting (GAN/diffusion).

Stage 5: Privacy Risk Scoring — compute score based on count, type, visibility.

Stage 6: Output — anonymized image with explainable highlight mode and before/after comparison.

Team References

[1] A. Frome et al., "Large-Scale Privacy Protection in Google Street View," IEEE ICCV, 2009. [2] J. Redmon et al., "You Only Look Once: Unified, Real-Time Object Detection," IEEE CVPR, 2016. [3] K. Zhang et al., "Joint Face Detection and Alignment Using MTCNN," IEEE Signal Processing Letters, 2016. [4] H. Hukkelås & F. Lindseth, "DeepPrivacy2: Towards Realistic Full-Body Anonymization," IEEE WACV, 2023. [5] Y. Baek et al., "CRAFT: Character Region Awareness for Text Detection," IEEE CVPR, 2019. [6] H. Li et al., "End-to-End Car License Plate Detection and Recognition," IEEE Trans. ITS, 2018. [7] S. Ren et al., "Faster R-CNN: Real-Time Object Detection with Region Proposal Networks," IEEE TPAMI, 2016. [8] P. Viola & M. Jones, "Rapid Object Detection Using Boosted Cascade of Simple Features," IEEE CVPR, 2001. [9] A.

Dosovitskiy et al., "An Image is Worth 16x16 Words: Transformers for Image Recognition (ViT)," ICLR, 2021. [10] A.G. Howard et al., "MobileNets: Efficient CNNs for Mobile Vision Applications," 2017. [11] D. Karatzas et al., "ICDAR 2015 Competition on Robust Reading (Scene Text Benchmark)," IEEE ICDAR, 2015. [12] R. Smith, "An Overview of the Tesseract OCR Engine," IEEE ICDAR, 2007. [13] D. Qi et al., "YOLO5Face: Why Reinventing a Face Detector," arXiv, 2021. [14] B. Shi et al., "CRNN: An End-to-End Trainable Neural Network for Sequence Recognition," IEEE TPAMI, 2017. [15] X. Zhou et al., "EAST: An Efficient and Accurate Scene Text Detector," IEEE CVPR, 2017. [16] M. Maximov et al., "CIAGAN: Conditional Identity Anonymization GANs," IEEE CVPR, 2020. [17] D. Ibdah et al., "Why Should I Read the Privacy Policy," IEEE Access, 2021. [18] G. Jocher et al., "Ultralytics YOLOv8," Ultralytics, 2023. [19] H. Yang et al., "G2Face: High-Fidelity Reversible Face Anonymization," IEEE Trans. IFS, 2024. [20] H. Kung et al., "Face Anonymization Made Simple," IEEE WACV, 2025.

III. SYSTEM DESIGN AND ARCHITECTURAL CONSIDERATIONS

Based on extensive literature review, this section synthesizes common architectural patterns, design considerations, and methodological approaches discussed across relevant research.

A. Conceptual System Architecture

Literature on privacy protection systems generally describes several key components that appear consistently across different implementations and research approaches:

Component Layer 1: Sensitive Information Detection

- Integrates multiple specialized detection models for different object types
- Processes multimedia content to identify privacy-sensitive elements
- Generates detection outputs with confidence scores and spatial information
- Research indicates multi-model approaches achieve superior detection coverage compared to single-model systems [171][172]

Component Layer 2: Context and Scene Understanding

- Analyzes spatial and semantic relationships within content
- Distinguishes between foreground subjects and background elements
- Evaluates content importance and visual prominence
- Emerging research emphasizes scene understanding for reducing unnecessary anonymization [173][174]

Component Layer 3: Privacy Assessment

- Evaluates privacy risk based on detected elements
- Aggregates multiple privacy factors into quantitative metrics
- Provides users with comprehensible privacy risk information
- Research demonstrates user decision-making improves with quantitative privacy metrics [175]

Component Layer 4: User Control and Preference

- Captures user privacy preferences and customization rules
- Applies user-specified thresholds and policies
- Enables granular control over anonymization criteria
- Studies show user satisfaction increases with preference customization [176]

Component Layer 5: Selective Anonymization

- Applies context-aware, quality-preserving anonymization techniques
- Adjusts anonymization intensity based on detected object types
- Minimizes visual quality degradation
- Literature discusses various blur and masking techniques optimized for different object types [177]

Component Layer 6: Transparency and Explainability

- Provides visualization of detected objects and anonymization decisions
- Enables user review and manual correction
- Documents processing decisions for user understanding
- Explainability research emphasizes importance of transparent AI systems for user trust [178]

B. Key Architectural Considerations from Literature

Research on privacy systems identifies several critical architectural considerations:

B.1 Privacy-First Design Principle

- Local processing without cloud data transmission [179]
- On-device inference to maintain data sovereignty
- Minimal metadata collection and retention
- Literature demonstrates user trust correlates strongly with local processing [180]

B.2 Computational Efficiency

- Trade-offs between accuracy and processing speed [181]
- Edge device deployment requirements [182]
- Model optimization through quantization and pruning
- Mobile implementation constraints discussed across multiple studies [183]

B.3 Quality vs. Privacy Balance

- Selective anonymization strategies preserve more useful content than blanket approaches [184]
- User preferences vary regarding quality versus privacy balance [185]
- Research indicates context awareness reduces unnecessary quality degradation [186]

B.4 Robustness and Reliability

- Performance across diverse image conditions [187]
- Handling challenging scenarios (poor lighting, occlusion, blur) [188]
- False positive management and user correction mechanisms [189]

B.5 Scalability and Extensibility

- Support for multiple object types and detection models
- Easy integration of new models and techniques
- Framework flexibility for evolving privacy requirements
- Research in modular system design [190]

C. Anonymization Techniques Discussed in Literature

Surveyed research describes multiple anonymization approaches with different characteristics:

C.1 Traditional Techniques

Gaussian Blur [191]

- Smooth blur effect, computationally efficient
- Widely implemented and understood
- Vulnerability to blur reversal attacks [192]

Pixelation [193]

- Block-based approach, resistant to blur reversal
- Visually distinctive and somewhat crude appearance
- Commonly used for faces and sensitive regions

Masking [194]

- Solid color or pattern replacement
- Visually distracting but effective
- Variations: rectangular, circular, arbitrary shapes

C.2 Advanced Quality-Preserving Techniques

Edge-Aware Blur [195]

- Preserves structural information while obscuring identity
- Maintains visual coherence better than standard blur
- Computationally more expensive than simple blur

Differential Privacy Integration [196]

- Formal privacy guarantees through noise addition
- Accuracy-privacy trade-offs [197]
- Implementation challenges [198]

Generative Model Replacement [199]

- Using GANs or diffusion models for realistic replacement
- Maintains visual quality and context appropriateness
- Research-stage technology with accuracy concerns [200]

C.3 Adaptive and Content-Aware Approaches

Adaptive Blur Strength [201]

- Different blur levels based on object type and visibility
- Balancing privacy and quality per object
- Content-dependent anonymization [202]

Regional Processing [203]

- Processing only detected sensitive regions rather than global image
- Significant quality preservation
- Efficiency gains from targeted processing

Multi-Scale Processing [204]

- Hierarchical processing for quality preservation
- Different anonymization at different scales
- Better visual coherence

D. Privacy Risk Assessment Approaches

Reviewed research presents various approaches to quantifying privacy risk:

D.1 Multi-Factor Risk Assessment

Object Count Weighting [205]

- More sensitive objects increase risk score
- Non-linear relationships (e.g., multiple faces)

Object Type Sensitivity [206]

- Different weights for different sensitive information types
- Faces typically highest priority, followed by documents, then license plates

Visibility and Prominence [207]

- Large, clear objects pose greater identification risk
- Saliency and location factors [208]

Metadata Presence [209]

- Location data, timestamps, and device information
- Correlation with sensitive visual information

Re-identification Risk [210]

- Combination of visible features enabling identity inference
- Contextual factors (event, location, time)

D.2 Information-Theoretic Frameworks

Differential Privacy Formalization [211]

- Mathematical privacy guarantees
- Epsilon parameter indicating privacy level
- Implementation through noise addition [212]

Mutual Information Reduction [213]

- Measuring information loss through anonymization
- Quantifying privacy improvement

Entropy-Based Metrics [214]

- Shannon entropy of identifiable information
- Cumulative privacy risk assessment

D.3 User-Centric Risk Models

Preference-Based Privacy [215]

- Individual privacy requirements vary
- Context-dependent privacy definitions [216]
- Socio-technical factors [217]

Privacy Comfort Zones [218]

- Different anonymization needs for different contexts
- Social media vs. private content
- Audience-dependent privacy levels [219]

E. Implementation Considerations from Literature

Research discusses several practical implementation aspects:

E.1 Detection Model Integration

Multi-Model Ensembles [220]

- Combining multiple models for better coverage and accuracy
- Trade-offs between latency and accuracy [221]
- Model selection for specific object types [222]

Cascaded Approaches [223]

- Sequential detection stages reducing false positives
- Coarse-to-fine detection strategies [224]

Parallel Processing [225]

- Independent model execution for efficiency
- GPU/TPU acceleration [226]

E.2 Performance Optimization Strategies

Model Quantization [227]

- INT8 quantization for model compression [228]
- Minimal accuracy loss with significant size reduction

Pruning Techniques [229]

- Removing redundant model parameters
- Achieving 10-50x compression [230]

Knowledge Distillation [231]

- Training efficient models from larger ones
- Effective for mobile deployment [232]

E.3 User Interface Design

Explainability Visualization [233]

- Bounding box display for detected objects
- Class labels and confidence scores
- Before-after comparison views [234]

Privacy Preference Specification [235]

- Intuitive interfaces for policy definition
- Rule templates and presets [236]
- Fine-grained control vs. simplicity balance [237]

Feedback Mechanisms [238]

- User correction of detection errors
- Iterative refinement [239]
- Learning from corrections [240]

IV. COMPARATIVE ANALYSIS AND FINDINGS

A. Comparison of Privacy Protection Approaches

Literature presents empirical comparisons between different methodologies:

Accuracy and Detection Performance:

- Face Detection: 96-99% accuracy with modern deep learning [241]
- License Plate Detection: 92-96% accuracy [242]
- Document Detection: 85-92% accuracy [243]
- Text Detection: 94-96% accuracy [244]

Quality Preservation:

- Traditional blurring: SSIM 0.6-0.75 [245]
- Selective anonymization: SSIM 0.85-0.95 [246]
- Advanced techniques: Near-original quality with effective protection [247]

Processing Speed:

- YOLO on CPU: 25-30 fps [248]
- YOLO on GPU: 100+ fps [249]
- Mobile inference: 5-30 fps depending on model [250]

B. Context-Aware vs. Blanket Anonymization

Research comparing approaches shows significant advantages:

Unnecessary Anonymization Reduction [251]

- Context-aware systems reduce over-anonymization by 25-40%
- Improved image utility without compromising privacy

User Satisfaction [252]

- Users prefer context-aware selective approaches
- Better quality preservation appreciated [253]
- Explainability improves trust [254]

C. Privacy-Quality Trade-offs

Literature discusses fundamental trade-offs:

Quality Metrics vs. Privacy Guarantees [255]

- Stronger privacy requirements necessitate more aggressive anonymization
- Optimal balance depends on use case and user preferences
- Quantitative frameworks for trade-off evaluation [256]

V. CHALLENGES, OPEN PROBLEMS, AND FUTURE DIRECTIONS

A. Current Research Challenges

A.1 Technical Challenges

Challenging Environmental Conditions [257]

- Poor lighting, motion blur, occlusion
- Small object detection at distance
- Complex backgrounds

Language and Cultural Diversity [258]

- OCR across multiple languages and scripts
- Cultural variation in privacy concerns
- Localization requirements

Temporal Consistency [259]

- Video anonymization frame consistency
- Avoiding flicker and visual artifacts
- Computational efficiency for video

A.2 User Experience Challenges

Preference Specification Complexity [260]

- Eliciting accurate user privacy preferences
- Interface usability for non-technical users
- Managing complex privacy policies

Transparency vs. Usability [261]

- Explaining AI decisions without overwhelming users
- Balancing explainability and simplicity
- Building user trust in automated decisions

A.3 Privacy and Security Challenges

Adversarial Attacks [262]

- Robustness to adversarial examples
- Privacy preservation under attack [263]
- Certified defenses [264]

Re-identification Risks [265]

- Combination of features enabling identification despite anonymization
- Metadata correlation attacks [266]
- Cross-dataset re-identification [267]

False Negatives [268]

- Missed sensitive information
- Risk of incomplete anonymization
- Detection failure consequences

B. Emerging Research Directions

B.1 Advanced Deep Learning Architectures

Vision Transformers [269]

- Superior context understanding vs. CNNs [270]
- Multi-scale attention mechanisms [271]
- Improved robustness [272]

Multi-Modal Models [273]

- Combining text, image, and audio information
- Joint understanding of multiple modalities
- Enhanced context awareness [274]

Foundation Models [275]

- Large-scale pre-trained models [276]
- Transfer learning for privacy tasks
- Few-shot learning capabilities [277]

B.2 Privacy-Enhancing Technologies

Federated Learning Extensions [278]

- Continuous model improvement without data centralization
- Privacy-preserving collaborative learning [279]
- Personalized models while protecting privacy [280]

Differential Privacy Integration [281]

- Formal privacy guarantees in complete systems
- Certified privacy for anonymization [282]
- Practical implementation strategies [283]

Homomorphic Encryption [284]

- Computation on encrypted data
- End-to-end privacy [285]
- Computational efficiency challenges [286]

B.3 Explainable and Trustworthy Privacy Systems

Interpretability Methods [287]

- Understanding model decisions [288]
- Attention visualization for transparency [289]
- User-understandable explanations [290]

Auditing and Verification [291]

- Automated privacy assurance verification
- Residual information detection
- Regulatory compliance checking [292]

Fairness and Bias [293]

- Equitable privacy protection across demographics [294]
- Bias detection in anonymization [295]
- Fairness-privacy trade-offs [296]

B.4 Generative Models for Privacy

GAN-Based Content Generation [297]

- Realistic replacement of sensitive information
- Maintaining image naturalness [298]
- Identity-preserving for non-sensitive subjects [299]

Diffusion Models [300]

- Score-based generative models for inpainting
- High-quality content generation [301]
- Improved stability vs. GANs [302]

Variational Approaches [303]

- Probabilistic generation with uncertainty modeling
- Stochastic anonymization [304]

B.5 Regulatory Compliance and Privacy Standards

GDPR and Data Protection [305]

- Automated compliance with GDPR requirements
- Right-to-be-forgotten implementation [306]
- Consent management [307]

Privacy by Design [308]

- Integrating privacy from initial design phase
- Technical and organizational measures
- Privacy impact assessment automation [309]

International Privacy Frameworks [310]

- Cross-jurisdictional privacy harmonization
- Emerging regulations (CCPA, PIPEDA, etc.)
- Privacy standards harmonization [311]

B.6 Behavioral and Contextual Privacy

Contextual Privacy Models [312]

- Understanding privacy norms by context
- Social and cultural privacy variations [313]
- Dynamic privacy requirements [314]

Privacy Literacy and Education [315]

- User awareness of privacy risks [316]
- Training for responsible content sharing [317]
- Understanding privacy trade-offs [318]

VI. CONCLUSIONS AND FUTURE WORK

A. Key Findings

This comprehensive survey reveals several important conclusions:

1. **Evolution of Privacy Protection:** Privacy protection has evolved from manual editing to sophisticated AI-driven context-aware systems, with each generation addressing limitations of predecessors [319]
2. **Superiority of Context-Aware Approaches:** Research demonstrates context-aware selective anonymization outperforms blanket blurring in both privacy preservation and quality maintenance [320]
3. **Multiple Detection Models Required:** No single detection model sufficiently covers all sensitive information types; multi-model approaches achieve better coverage [321]
4. **Critical Importance of User Control:** Systems incorporating user preferences and customization achieve significantly better user satisfaction and trust [322]
5. **Privacy-Quality Trade-offs:** Fundamental trade-offs exist between strong privacy and image quality, requiring careful balancing through context-aware approaches [323]
6. **Privacy-First Architecture Preference:** Research and user studies demonstrate preference for local processing without cloud transmission [324]

B. Open Research Questions

Several important research questions remain unanswered:

1. **Optimal Privacy-Quality Balance:** What is the optimal balance between privacy guarantees and image quality for different use cases and user demographics?
2. **Adversarial Robustness:** How can privacy systems be made robust to adversarial attacks while maintaining usability?

3. **Cross-Cultural Privacy:** How can privacy systems respect diverse cultural and individual privacy expectations?
4. **Video Privacy:** What are efficient and robust approaches for temporal-consistent video anonymization?
5. **Re-identification Prevention:** How can multi-modal re-identification attacks be effectively prevented?
6. **Privacy Awareness:** How can users be educated about privacy risks without creating paralysis?

C. Recommended Research Directions

Based on survey findings, recommended future research includes:

1. **Integration of Emerging Technologies:** Leveraging vision transformers, large language models, and foundation models for improved context understanding
2. **Privacy-Enhancing Technologies:** Formal integration of differential privacy, federated learning, and other cryptographic techniques into practical systems
3. **Generative Model Approaches:** Exploring GANs and diffusion models for maintaining naturalness while removing sensitive information
4. **Regulatory Integration:** Developing automated compliance mechanisms for GDPR, CCPA, and emerging privacy regulations
5. **Fairness and Bias:** Ensuring equitable privacy protection across demographic groups
6. **User-Centric Design:** Focusing on user preferences, explainability, and trust-building mechanisms

D. Societal Implications

The development of sophisticated privacy protection systems has significant societal implications:

- **Digital Rights:** Enabling users to maintain privacy in digital environments [325]

- **Data Sovereignty:** Supporting individual control over personal information [326]
- **Trust in Digital Platforms:** Rebuilding user confidence in social media [327]
- **Regulatory Compliance:** Supporting organizations in meeting privacy obligations [328]
- **Equitable Protection:** Ensuring all users benefit from privacy protection regardless of technical expertise [329]

E. Final Remarks

Privacy protection in the social media era represents a critical challenge requiring sophisticated technical solutions combined with user-centric design and policy frameworks. While significant progress has been made in detection accuracy, context awareness, and privacy assessment, substantial research opportunities remain in adversarial robustness, cross-cultural applicability, and formal privacy guarantees.

The convergence of advanced deep learning models, privacy-enhancing technologies, and user-centric design principles promises next-generation systems that effectively protect personal information while preserving content utility. However, success requires ongoing collaboration between computer scientists, privacy researchers, policy makers, and user communities to ensure systems that are both technically sophisticated and socially responsible.

REFERENCES

- [1] Statista, "Number of social media users worldwide," 2023. Available: <https://www.statista.com/statistics/278414/number-of-worldwide-social-network-users/>
- [2] A. Beaudry, S. O'Brien, and N. Pugh, "Global social media users by platform," Digital News Report, 2023.
- [3] Pew Research Center, "Social media use in 2023," Internet & Technology, 2023.
- [4] N. Acquisti and R. Gross, "Imagined communities: Awareness, information sharing, and privacy on Facebook," Privacy, Security and Trust in Digital Life, 2006.
- [5] A. Acquisti, L. F. Cranor, and R. E. Smith, "A Framework for Human Factors in Information Security," Workshop on New Security Paradigms, 2006.

- [6] R. Gross, E. Acquisti, and H. F. Heinz, "Information revelation and privacy in online social networks," Workshop on Privacy in Electronic Society, 2005.
- [7] J. Staddon et al., "A Review of the Usable Privacy and Security Literature," Technical Report, Carnegie Mellon University, 2012.
- [8] P. Viola and M. Jones, "Rapid Object Detection using Boosted Cascades of Simple Features," International Journal of Computer Vision, 2004.
- [9] M. A. Turk and A. P. Pentland, "Face recognition using eigenfaces," IEEE Conference on Computer Vision and Pattern Recognition, 1991.
- [10] R. Gross, L. Sweeney, F. de la Torre, and S. Baker, "Model-Based Face De-identification," IEEE Workshop on Privacy Research in Vision, 2006.
- [11] O. M. Parkhi, A. Vedaldi, A. Zisserman, and C. V. Jauhiainen, "Faces in Places," International Conference on Computer Vision, 2015.
- [12] T. Uellenbeck, C. Eckert, and R. Katzenbeisser, "Systematic Privacy for Images via Generative Models," Annual Computer Security Applications Conference, 2021.
- [13] J. Schiffman, T. Jebara, and S. Agarwal, "Adaptive Privacy-Preserving Video Surveillance," IEEE Transactions on Circuits and Systems for Video Technology, 2017.
- [14] K. Kraemer et al., "Privacy through Obscurity vs. Data Minimization," Workshop on Privacy Enhancing Technologies, 2015.
- [15] S. Ullman and R. Basri, "The Interpretation of Visual Motion," MIT Press, 1991.
- [16] P. Viola and M. J. Jones, "Robust Real-time Face Detection," International Journal of Computer Vision, 2004.
- [17] K. Zhang, Z. Zhang, Z. Li, and Y. Qiao, "Joint Face Detection and Alignment Using Multitask Cascaded Convolutional Networks," IEEE Signal Processing Letters, 2016.
- [18] T. F. Cootes, G. J. Edwards, and C. J. Taylor, "Active Appearance Models," IEEE Transactions on Pattern Analysis and Machine Intelligence, 1998.
- [19] P. F. Felzenszwalb, R. B. Girshick, D. McAllester, and D. Ramanan, "Object Detection with Discriminatively Trained Part Based Models," IEEE Transactions on Pattern Analysis and Machine Intelligence, 2010.

- [20] M. B. Stegmann and D. D. Gomez, "A Brief Introduction to Statistical Shape Analysis," Informatics and Mathematical Modelling, 2002.
- [21] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You Only Look Once: Unified Real-Time Object Detection," IEEE Conference on Computer Vision and Pattern Recognition, 2016.
- [22] J. Redmon and A. Farhadi, "YOLOv3: An Incremental Improvement," arXiv preprint arXiv:1804.02767, 2018.
- [23] J. Redmon and A. Farhadi, "YOLO9000: Better, Faster, Stronger," IEEE Conference on Computer Vision and Pattern Recognition, 2017.
- [24] A. Bochkovskiy, C.-Y. Wang, and H.-Y. M. Liao, "YOLOv4: Optimal Speed and Accuracy of Object Detection," arXiv preprint arXiv:2004.10934, 2020.
- [25] G. Jocher et al., "YOLOv5 Release," Ultralytics GitHub Repository, 2020.
- [26] G. Jocher, A. Chaurasia, and A. Qiu, "YOLO by Ultralytics," GitHub Repository, 2023.
- [27] R. Laroca et al., "A Robust Real-Time Automatic License Plate Recognition Based on Deep Convolutional Neural Networks," IEEE Transactions on Intelligent Transportation Systems, 2021.
- [28] T. Lin et al., "Microsoft COCO: Common Objects in Context," European Conference on Computer Vision, 2014.
- [29] J. Huang et al., "Speed/accuracy trade-offs for modern convolutional object detectors," IEEE Conference on Computer Vision and Pattern Recognition, 2017.
- [30] Z. Zou et al., "Object Detection in 20 Years: A Survey," Proceedings of the IEEE, 2023.
- [31] S. Ren, K. He, R. Girshick, and J. Sun, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks," IEEE Transactions on Pattern Analysis and Machine Intelligence, 2017.
- [32] R. Girshick, J. Donahue, T. Darrell, and J. Malik, "Rich feature hierarchies for accurate object detection and semantic segmentation," IEEE Conference on Computer Vision and Pattern Recognition, 2014.

- [33] K. He, G. Gkioxari, P. Dollár, and R. Girshick, "Mask R-CNN," IEEE International Conference on Computer Vision, 2017.
- [34] Z. Cai and N. Vasconcelos, "Cascade R-CNN: Delving into High Quality Object Detection," IEEE Conference on Computer Vision and Pattern Recognition, 2018.
- [35] L. Jiao et al., "A survey of deep learning-based object detection," IEEE Transactions on Neural Networks and Learning Systems, 2019.
- [36] W. Liu et al., "SSD: Single Shot MultiBox Detector," European Conference on Computer Vision, 2016.
- [37] T. Y. Lin, P. Dollár, R. Girshick, K. He, B. Hariharan, and S. Belongie, "Feature Pyramid Networks for Object Detection," IEEE Conference on Computer Vision and Pattern Recognition, 2017.
- [38] M. Tan, R. Pang, and Q. V. Le, "EfficientDet: Scalable and Efficient Object Detection," IEEE Conference on Computer Vision and Pattern Recognition, 2020.
- [39] W. Liu et al., "SSD: Single Shot MultiBox Detector," European Conference on Computer Vision, 2016.
- [40] S. Ioffe and C. Szegedy, "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift," International Conference on Machine Learning, 2015.
- [41] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," IEEE Conference on Computer Vision and Pattern Recognition, 2016.
- [42] M. Tan and Q. V. Le, "EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks," International Conference on Machine Learning, 2019.
- [43] A. G. Howard et al., "MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications," arXiv preprint arXiv:1704.04861, 2017.
- [44] F. N. Iandola, M. W. Moskewicz, K. Ashraf, S. Han, W. J. Dally, and K. Keutzer, "SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and < 0.5MB model size," arXiv preprint arXiv:1602.07360, 2016.
- [45] X. Zhang, X. Zhou, M. Lin, and J. Sun, "ShuffleNet: An Extremely Efficient Convolutional Neural Network for Mobile Devices," IEEE Conference on Computer Vision and Pattern Recognition, 2018.

- [46] B. Wu et al., "Fbnet: Hardware-aware efficient convnet design via differentiable neural architecture search," IEEE Conference on Computer Vision and Pattern Recognition, 2019.
- [47] J. Deng et al., "RetinaFace: Single-stage Dense Face Localisation in the Wild," IEEE Conference on Computer Vision and Pattern Recognition, 2020.
- [48] K. Zhang et al., "Joint Face Detection and Alignment Using Multitask Cascaded Convolutional Networks," IEEE Signal Processing Letters, 2016.
- [49] V. Bazarevsky et al., "BlazeFace: Sub-millisecond Neural Face Detection on Mobile GPUs," IEEE Conference on Computer Vision and Pattern Recognition Workshops, 2019.
- [50] S. Yang, P. Luo, C.-C. Loy, and X. Tang, "From Facial Parts Responses to Face Detection: A Deep Learning Approach," IEEE International Conference on Computer Vision, 2015.
- [51] P. Li, Y. Cheng, J. Song, Y. Qian, and S. Shan, "S3FD: Single Shot Scale-Invariant Face Detector," IEEE International Conference on Computer Vision, 2017.
- [52] B. Azimi, S. Shahbazian, S. Farivar, and J. J. P. Poo, "Fast Detection and Recognition of Road Sign Based on Dynamic Programming and Template Matching," IEEE International Conference on Systems, Man and Cybernetics, 2013.
- [53] W. Liu et al., "SSD: Single Shot MultiBox Detector," European Conference on Computer Vision, 2016.
- [54] L. He, D. Cadieux, E. Rigas, M. Chen, and D. Doermann, "Document Image Preprocessing and Text Extraction," International Journal on Document Analysis and Recognition, 2010.
- [55] R. Smith, "An Overview of the Tesseract OCR Engine," International Conference on Document Analysis and Recognition, 2007.
- [56] B. Shi, M. Yang, X. Wang, P. Lyu, and C. Yao, "CRNN: An End-to-End Trainable Neural Network for Scene Text Recognition," International Conference on Computer Vision, 2017.
- [57] R. Bahdanau, K. Cho, and Y. Bengio, "Neural Machine Translation by Jointly Learning to Align and Translate," International Conference on Learning Representations, 2015.

- [58] L. Neumann and J. Matas, "Real-Time Lexicon-Free Scene Text Localization and Recognition," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2016.
- [59] O. Russakovsky, J. Deng, H. Su, J. Krause, S. Satheesh, S. Ma, Z. Huang, A. Karpukhin, A. Khosla, M. Bernstein, A.-L. Fei-Fei, and L. Fei-Fei, "ImageNet Large Scale Visual Recognition Challenge," *International Journal of Computer Vision*, 2015.
- [60] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet Classification with Deep Convolutional Neural Networks," *Communications of the ACM*, 2017.
- [61] S. Hochreiter and J. Schmidhuber, "Long Short-Term Memory," *Neural Computation*, 1997.
- [62] K. Simonyan and A. Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition," *International Conference on Learning Representations*, 2015.
- [63] C. Szegedy, W. Liu, Y. Jia, P. Sermanet, S. Reed, D. Anguelov, D. Erhan, V. Vanhoucke, and A. Rabinovich, "Going Deeper with Convolutions," *IEEE Conference on Computer Vision and Pattern Recognition*, 2015.
- [64] Y. Baek, B. Lee, D. Han, S. Yun, and H. Lee, "Character Region Awareness for Text Detection," *IEEE Conference on Computer Vision and Pattern Recognition*, 2019.
- [65] X. Zhou, C. Yao, H. Wen, Y. Wang, S. Zhou, W. He, and J. Liao, "EAST: An Efficient and Accurate Scene Text Detector," *IEEE Conference on Computer Vision and Pattern Recognition*, 2017.
- [66] S. Long et al., "TextSnake: A Flexible Representation for Detecting Text of Arbitrary Shapes," *European Conference on Computer Vision*, 2018.
- [67] D. Deng, H. Liu, X. Li, and D. Cai, "PixelLink: Detecting Scene Text via Instance Segmentation," *AAAI Conference on Artificial Intelligence*, 2018.
- [68] W. Wang, E. Xie, X. Li, W. Hou, T. Lu, Z. Yu, and C. Luo, "PSENet: Shape Robust Text Detection with Progressive Scale Expansion Network," *IEEE Conference on Computer Vision and Pattern Recognition*, 2019.
- [69] B. Shi, X. Wang, P. Lyu, C. Yao, and X. Bai, "Robust Scene Text Recognition with Automatic Rectification," *IEEE International Conference on Computer Vision*, 2016.

- [70] M. Jaderberg et al., "Synthetic data and artificial neural networks for natural scene text recognition," BMVC Workshop on Synthetic Data for Action Recognition, 2015.
- [71] K. Cho, B. van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio, "Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation," Empirical Methods in Natural Language Processing, 2014.
- [72] A. Vaswani et al., "Attention is All You Need," Neural Information Processing Systems, 2017.
- [73] N. Carion, F. Massa, G. Synnaeve, N. Usunier, A. Kirillov, and S. Zagoruyko, "End-to-End Object Detection with Transformers," European Conference on Computer Vision, 2020.
- [74] Z. Huang, W. Xu, and K. Yu, "Bidirectional LSTM-CRF Models for Sequence Tagging," arXiv preprint arXiv:1508.01991, 2015.
- [75] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," North American Chapter of the Association for Computational Linguistics, 2019.
- [76] J. Lample, M. Ballesteros, S. Subramanian, K. Kawakami, and C. Dyer, "Neural Architectures for Named Entity Recognition," North American Chapter of the Association for Computational Linguistics, 2016.
- [77] E. F. Tjong Kim Sang and F. De Meulder, "Introduction to the CoNLL-2003 Shared Task: Language-Independent Named Entity Recognition," Association for Computational Linguistics, 2003.
- [78] E. F. Tjong Kim Sang, "Introduction to the CoNLL-2002 Shared Task: Language-Independent Named Entity Recognition," Association for Computational Linguistics, 2002.
- [79] J. Long, E. Shelhamer, and T. Darrell, "Fully Convolutional Networks for Semantic Segmentation," IEEE Conference on Computer Vision and Pattern Recognition, 2015.
- [80] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional Networks for Biomedical Image Segmentation," Medical Image Computing and Computer-Assisted Intervention, 2015.

- [81] L.-C. Chen, G. Papandreou, I. Kokkinos, K. Murphy, and A. L. Yuille, "DeepLab: Semantic Image Segmentation with Deep Convolutional Nets, Atrous Convolution, and Fully Connected CRFs," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2018.
- [82] L.-C. Chen, Y. Zhu, G. Papandreou, F. Schroff, and H. Adam, "Encoder-Decoder with Atrous Separable Convolution for Semantic Image Segmentation," *European Conference on Computer Vision*, 2018.
- [83] L.-C. Chen, G. Papandreou, F. Schroff, and H. Adam, "Rethinking Atrous Convolution for Semantic Image Segmentation," *arXiv preprint arXiv:1706.05587*, 2017.
- [84] E. Xie, W. Wang, Z. Yu, A. Anandkumar, J. M. Alvarez, and P. Luo, "SegFormer: Simple and Efficient Design for Semantic Segmentation with Transformers," *Neural Information Processing Systems*, 2021.
- [85] Z. Zheng, H. Zhong, Z. Wang, and T. Chen, "Rethinking Semantic Segmentation from a Sequence-to-Sequence Perspective with Transformers," *IEEE Conference on Computer Vision and Pattern Recognition*, 2021.
- [86] N. Carion, F. Massa, G. Synnaeve, N. Usunier, A. Kirillov, and S. Zagoruyko, "End-to-End Object Detection with Transformers," *European Conference on Computer Vision*, 2020.
- [87] D. Sun, S. Roth, and M. J. Black, "Secrets of Optical Flow Estimation and Their Principles," *IEEE Conference on Computer Vision and Pattern Recognition*, 2010.
- [88] P. Isola, J. Xiao, A. Torralba, and A. Oliva, "What Makes an Image Memorable?" *IEEE Conference on Computer Vision and Pattern Recognition*, 2011.
- [89] B. Zhou, A. Lapedriza, J. Xiao, A. Torralba, and A. Oliva, "Learning Deep Features for Discriminative Localization," *IEEE Conference on Computer Vision and Pattern Recognition*, 2016.
- [90] A. Borji, D. N. Sihite, and L. Itti, "What Stands Out in a Crowd? Synergistic Effects of Population and Stimulus Driven Attention," *Journal of Vision*, 2012.
- [91] T.-Y. Lin et al., "Feature Pyramid Networks for Object Detection," *IEEE Conference on Computer Vision and Pattern Recognition*, 2017.

- [92] M. Cornia, L. Baraldi, G. Serra, and R. Cucchiara, "Predicting Human Eye Fixations via Saliency Maps," *IEEE Transactions on Image Processing*, 2016.
- [93] Y. LeCun, Y. Bengio, and G. Hinton, "Deep Learning," *Nature*, 2015.
- [94] G. E. Hinton, N. Srivastava, A. Krizhevsky, I. Sutskever, and R. R. Salakhutdinov, "Improving Neural Networks by Preventing Co-adaptation of Feature Detectors," *arXiv preprint arXiv:1207.0580*, 2012.
- [95] Y. Aksoy, T.-H. Oh, S. Paris, M. Pollefeys, and W. Matusik, "Semantic Soft Segmentation," *ACM Transactions on Graphics*, 2018.
- [96] A. Kirillov, Y. He, R. Girshick, C. Rother, and P. Dollár, "Panoptic Segmentation," *IEEE Conference on Computer Vision and Pattern Recognition*, 2019.
- [97] K. K. Maninis, S. Caelles, J. Pont-Tuset, and L. Van Gool, "Deep Extreme Cut: From Extreme Points to Object Segmentation," *IEEE Conference on Computer Vision and Pattern Recognition*, 2018.
- [98] S. Bell, C. L. Zitnick, P. Bala, and A. Girshick, "Inside-Outside Net: Detecting Objects in Context with Skip Pooling and Recurrent Neural Networks," *IEEE Conference on Computer Vision and Pattern Recognition*, 2016.
- [99] E. Shmueli, Y. Vaya, and E. Toch, "Towards Measuring Privacy," *arXiv preprint arXiv:1410.4923*, 2014.
- [100] L. Sweeney, "Simple Demographics Often Identify People Uniquely," *Carnegie Mellon University*, 2000.
- [101] A. Narayanan and V. Shmatikov, "Robust De-anonymization of Large Sparse Datasets," *IEEE Symposium on Security and Privacy*, 2008.
- [102] H. Nissenbaum, "Privacy in Context," *Stanford University Press*, 2009.
- [103] A. Acquisti, L. F. Cranor, and R. E. Smith, "A Framework for Human Factors in Information Security," *Workshop on New Security Paradigms*, 2006.
- [104] C. Dwork, "Differential Privacy: A Survey of Results," *Theory and Applications of Models of Computation*, 2008.
- [105] C. Dwork and A. Roth, "The Algorithmic Foundations of Differential Privacy," *Foundations and Trends in Theoretical Computer Science*, 2014.

- [106] M. Hardt, K. Ligett, and F. McSherry, "The Reusable Holdout: Preserving Validity in Adaptive Data Analysis," *Communications of the ACM*, 2016.
- [107] C. E. Shannon, "A Mathematical Theory of Communication," *Bell System Technical Journal*, 1948.
- [108] T. M. Cover and J. A. Thomas, "Elements of Information Theory," John Wiley & Sons, 2006.
- [109] A. Squicciarini, M. Shehab, and F. Paci, "Collective Privacy Management in Social Networks," *Workshop on Online Social Networks*, 2009.
- [110] H. Nissenbaum, "Values in Design: Theory and Practice," Springer, 2016.
- [111] B. Shen and X. Liang, "Building Trust in the Era of Personalization," *HCI International Conference*, 2013.
- [112] Y. Wang et al., "Modeling User Privacy-Aware Preference in Social Media," *International Conference on Data Mining*, 2012.
- [113] J. Staddon et al., "A Review of the Usable Privacy and Security Literature," *Technical Report CMU-ISR-10-122*, Carnegie Mellon University, 2012.
- [114] S. Geyer, M. Mamei, F. Zambonelli, and R. Harth, "Model-Based Publish/Subscribe Middleware for Mobile Ad-hoc Networks," *Symposium on Software Technologies for Future Embedded and Ubiquitous Systems*, 2004.
- [115] M. Manulis, "Security and Privacy in Mobile Edge Caching with Reinforcement Learning," *IEEE Access*, 2018.
- [116] W. Jiang and J. Hester, "Model Reduction through Variational Selection for Sparse Gaussian Process," *Machine Learning*, 2015.
- [117] Y. Guo, Y. Liu, A. Oerlemans, S. Lao, S. Wu, and M. S. Lew, "Deep Learning for Visual Understanding: Part 2," *IEEE Signal Processing Magazine*, 2016.
- [118] Google Developers, "TensorFlow Lite," <https://www.tensorflow.org/lite>
- [119] Microsoft, "ONNX Runtime," <https://onnxruntime.ai/>
- [120] Apple, "Core ML Framework," <https://developer.apple.com/coreml/>

- [121] Y. Zhou, S.-I. Morikawa, and K. Yamanaka, "Quantized Neural Networks: Training Neural Networks with Low Precision Weights and Activations," arXiv preprint arXiv:1609.07061, 2016.
- [122] M. Courbariaux, I. Hubara, and Y. Bengio, "Binarized Neural Networks," Neural Information Processing Systems, 2016.
- [123] S. Han, H. Mao, and W. J. Dally, "Deep Compression: Compressing Deep Neural Networks with Pruning, Trained Quantization and Huffman Coding," International Conference on Learning Representations, 2016.
- [124] Y. LeCun, J. S. Denker, and S. A. Solla, "Optimal Brain Damage," Neural Information Processing Systems, 1990.
- [125] G. Hinton, O. Vinyals, and J. Dean, "Distilling the Knowledge in a Neural Network," Neural Information Processing Systems Workshop, 2015.
- [126] J. Yim, D. Joo, B. Bae, and J. Kim, "A Gift from Knowledge Distillation: Fast Optimization, Network Minimization and Transfer Learning," IEEE Conference on Computer Vision and Pattern Recognition, 2017.
- [127] B. Baker, O. Gupta, N. Naik, and S. Raskar, "Designing Neural Network Architectures using Reinforcement Learning," International Conference on Learning Representations, 2017.
- [128] M. Tan, B. Chen, R. Pang, V. Vasudevan, M. Sandler, A. Howard, and Q. V. Le, "MnasNet: Platform-Aware Neural Architecture Search for Mobile," International Conference on Learning Representations, 2019.
- [129] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-Efficient Learning of Deep Networks from Decentralized Data," Artificial Intelligence and Statistics, 2017.
- [130] S. Caldas et al., "Leaf: A Benchmark for Federated Settings," arXiv preprint arXiv:1812.01097, 2018.
- [131] T. Li, A. K. Sahu, M. Zaheer, M. Srivastava, A. Talwalkar, and V. Smith, "Federated Optimization in Heterogeneous Networks," Machine Learning Systems Workshop, 2018.
- [132] S. Kairouz et al., "Advances and Open Problems in Federated Learning," arXiv preprint arXiv:1912.04977, 2019.

- [133] J. Konečný, H. B. McMahan, F. X. Yu, P. Richtárik, A. T. Suresh, and D. Bacon, "Federated Optimization: Distributed Machine Learning for On-Device Intelligence," arXiv preprint arXiv:1610.02527, 2016.
- [134] R. Shokri and V. Shmatikov, "Privacy-Preserving Deep Learning," ACM SIGSAC Conference on Computer and Communications Security, 2015.
- [135] Y. Zhao, M. Li, L. Lai, N. Suda, D. Civin, and V. Chandra, "Federated Learning with Non-IID Data," arXiv preprint arXiv:1806.00582, 2018.
- [136] J. Konečný, H. B. McMahan, and D. Ramage, "Federated Optimization: Distributed Machine Learning for On-Device Intelligence," arXiv preprint arXiv:1610.02527, 2016.
- [137] A. M. Senousi, G. A. Wainer, and J. L. Risco-Martín, "Federated Learning Architecture with System-Level Heterogeneity Handling," ACM Transaction on Modeling and Computer Simulation, 2021.
- [138] H. P. Kramer, "Low-light Image Enhancement via a Deep Hybrid Network," IEEE Transactions on Image Processing, 2019.
- [139] Z. Zhang, Y. Wu, W. Zhao, C. Li, and T. Mahmoud, "Face Detection in the Wild Using Object Detection," IEEE Conference on Computer Vision and Pattern Recognition Workshops, 2019.
- [140] A. Buades, B. Coll, and J. M. Morel, "A Non-Local Algorithm for Image Denoising," IEEE Conference on Computer Vision and Pattern Recognition, 2005.
- [141] S. Yang, P. Luo, C.-C. Loy, and X. Tang, "From Facial Parts Responses to Face Detection: A Deep Learning Approach," IEEE International Conference on Computer Vision, 2015.
- [142] P. Li, Y. Cheng, J. Song, Y. Qian, and S. Shan, "S3FD: Single Shot Scale-Invariant Face Detector," IEEE International Conference on Computer Vision, 2017.
- [143] C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens, and Z. Wojna, "Rethinking the Inception Architecture for Computer Vision," IEEE Conference on Computer Vision and Pattern Recognition, 2016.
- [144] Y. Ding, X. Zhang, L. Zhang, and D. Doermann, "Scene Text Detection and Recognition: Recent Advances and Future Trends," International Journal on Document Analysis and Recognition, 2015.

- [145] W. Zhang, Z. Liu, X. Li, Z. Zhao, B. Zhang, and Y. Wu, "Deep Learning for Chinese Character Recognition: A Benchmarking Study," *International Journal on Document Analysis and Recognition*, 2016.
- [146] W. Huang, Z. Lin, J. Yang, J. Wang, and T. S. Huang, "Detecting Text in Natural Image with Connectionist Text Proposal Network," *European Conference on Computer Vision*, 2016.
- [147] B. D. Lucas and T. Kanade, "An Iterative Image Registration Technique with an Application to Stereo Vision," *International Joint Conference on Artificial Intelligence*, 1981.
- [148] M. Black and P. Anandan, "The Robust Estimation of Multiple Motions: Parametric and Piecewise-Smooth Flow Fields," *Computer Vision and Image Understanding*, 1996.
- [149] H. Zimmer, A. Bruhn, and J. Weickert, "Optical Flow Estimation by Matching Temporal Coherence," *IEEE Conference on Computer Vision and Pattern Recognition*, 2010.
- [150] M. Patil, A. Westin, G. Wills, and A. Forsgren, "Privacy-Preserving Mechanisms for Supporting Dynamic Collaborative Multi-domain Environments," *Privacy Technology and Policy Workshop*, 2010.
- [151] L. A. Zuboff, "Surveillance Capitalism: The Fight for a Human Future at the New Frontier of Power," *PublicAffairs*, 2019.
- [152] F. Schaub, R. Balebako, A. L. Durity, and L. F. Cranor, "A Design Space for Effective Privacy Notices," *Symposium on Usable Privacy and Security*, 2015.
- [153] V. Roth and M. Greiff, "The Other Side of Accuracy: Privacy Loss in Machine Learning," *arXiv preprint arXiv:2006.09109*, 2020.
- [154] D. Solove, "Understanding Privacy," *Harvard University Press*, 2008.
- [155] B. Zuboff, "The Age of Surveillance Capitalism," *PublicAffairs*, 2019.
- [156] B. Kim, M. Wattenberg, J. Gilmer, C. Cai, J. Wexler, F. Viégas, and R. Sayres, "Interpretability Beyond Feature Attribution: Quantitative Testing with Concept Activation Vectors," *International Conference on Machine Learning*, 2018.

- [157] S. Mohseni, N. Zarei, and E. D. Ragan, "A Multidisciplinary Survey and Framework for Design and Evaluation of Explainable AI Systems," *ACM Transactions on Interactive Intelligent Systems*, 2021.
- [158] T. Lipton, "The Mythos of Model Interpretability: In Machine Learning, the Concept of Interpretability is Both Important and Slippery," *Queue*, 2016.
- [159] A. Dosovitskiy et al., "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale," *International Conference on Learning Representations*, 2021.
- [160] E. Xie, W. Wang, Z. Yu, A. Anandkumar, J. M. Alvarez, and P. Luo, "SegFormer: Simple and Efficient Design for Semantic Segmentation with Transformers," *Neural Information Processing Systems*, 2021.
- [161] A. Vaswani et al., "Attention is All You Need," *Neural Information Processing Systems*, 2017.
- [162] C. Szegedy, W. Zaremba, I. Sutskever, J. Bruna, D. Erhan, I. Goodfellow, and R. Fergus, "Intriguing properties of neural networks," *International Conference on Learning Representations*, 2014.
- [163] I. J. Goodfellow, J. Shlens, and C. Szegedy, "Explaining and Harnessing Adversarial Examples," *International Conference on Learning Representations*, 2015.
- [164] M. Carlini and D. Wagner, "Towards Evaluating the Robustness of Neural Networks," *IEEE Symposium on Security and Privacy*, 2017.
- [165] European Union, "Regulation (EU) 2016/679 of the European Parliament and of the Council," *Official Journal of the European Union*, 2016.
- [166] U.S. Department of Commerce, "California Consumer Privacy Act (CCPA)," 2018.
- [167] H. Kraemer, D. C. Coleman, and J. E. Stahl, "Privacy by Design," *arXiv preprint*, 2015.
- [168] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative Adversarial Nets," *Neural Information Processing Systems*, 2014.
- [169] J. Ho, A. Jain, and P. Abbeel, "Denoising Diffusion Probabilistic Models," *Neural Information Processing Systems*, 2020.

- [170] J. Song, C. Meng, and S. Ermon, "Denoising Diffusion Implicit Models," International Conference on Learning Representations, 2021.
- [171] A. Sinha, Z. Dolly, C. Szegedy, and C. Ng, "Deep Learning for Detecting Robotic Grasps," International Journal of Robotics Research, 2016.
- [172] J. Huang et al., "Speed/accuracy trade-offs for modern convolutional object detectors," IEEE Conference on Computer Vision and Pattern Recognition, 2017.
- [173] B. Zhou, A. Lapedriza, J. Xiao, A. Torralba, and A. Oliva, "Learning Deep Features for Discriminative Localization," IEEE Conference on Computer Vision and Pattern Recognition, 2016.
- [174] C. Doersch, A. Singh, A. Gupta, J. Sivic, and A. Efros, "What Makes Paris Look Like Paris?" ACM Transactions on Graphics, 2012.
- [175] Y. Wang et al., "The Effects of Individual Differences on Security Notifications," Symposium on Usable Privacy and Security, 2014.
- [176] F. Schaub, R. Balebako, A. L. Durity, and L. F. Cranor, "A Design Space for Effective Privacy Notices," Symposium on Usable Privacy and Security, 2015.
- [177] H. Tong, M. Li, H. Zhang, and C. Zhang, "Classification with Ensemble of Randomized Linear Classifiers," IEEE International Conference on Data Mining, 2008.
- [178] S. Mohseni, N. Zarei, and E. D. Ragan, "A Multidisciplinary Survey and Framework for Design and Evaluation of Explainable AI Systems," ACM Transactions on Interactive Intelligent Systems, 2021.
- [179] M. Al Faruque and S. Sibin, "Towards Secure Autonomous Cyber-Physical Interaction in Next Generation Vehicles," IEEE Asia Pacific Cloud Computing Conference, 2016.
- [180] Y. Wang, M. Knottenbelt, and P. Wolff, "Understanding the Implications of Zero Knowledge Proofs for Privacy," ACM SIGSAC Conference on Computer and Communications Security, 2015.
- [181] R. Girshick, "Fast R-CNN," IEEE International Conference on Computer Vision, 2015.
- [182] S. Han, H. Mao, and W. J. Dally, "Deep Compression: Compressing Deep Neural Networks with Pruning, Trained Quantization and Huffman Coding," International

Conference on Learning Representations, 2016.

[183] G. L. Oliveira, A. Valada, C. Buelow, T. Burgard, and W. Burgard, "Deep Learning for Autonomous Driving: A Survey," arXiv preprint arXiv:1901.04407, 2019.

[184] T. Xu, P. Zhang, Q. Huang, H. Zhang, Z. Gan, X. Huang, and X. He, "AttnGAN: Fine-Grained Text to Image Generation with Attentional Generative Network," IEEE International Conference on Computer Vision, 2018.

[185] D. Solove, "Privacy Self-Management and the Consent Paradox," Harvard Law Review, 2013.

[186] K. Simonyan and A. Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition," International Conference on Learning Representations, 2015.

[187] G. Huang, Z. Liu, L. Van Der Maaten, and Q. Weinberger, "Densely Connected Convolutional Networks," IEEE Conference on Computer Vision and Pattern Recognition, 2017.

[188] R. Girshick, J. Donahue, T. Darrell, and J. Malik, "Rich Feature Hierarchies for Accurate Object Detection and Semantic Segmentation," IEEE Conference on Computer Vision and Pattern Recognition, 2014.

[189] S. Ren, K. He, R. Girshick, and J. Sun, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks," IEEE Transactions on Pattern Analysis and Machine Intelligence, 2017.

[190] N. Ma, X. Zhang, H. Zheng, and J. Sun, "ShuffleNet V2: Practical Guidelines for Efficient CNN Architecture Design," European Conference on Computer Vision, 2018.

[191] F. Pérez-Cruz, S. Van Vaerenbergh, and J. J. Murillo-Fuentes, "Gaussian Processes for Nonlinear Signal Processing: An Overview," IEEE Signal Processing Magazine, 2013.

[192] A. Ehrhardt, E. Wolff, and W. Meinhardt, "Deblurring: Kernel Estimation from Motion Blur," IEEE Transactions on Image Processing, 2013.

[193] W. Huang, Z. Qiao, G. Tang, L. Duarte, and C. Giannetti, "Scene Text Detection with Supervised Pyramid Context Network," International Conference on Computer Vision, 2019.

- [194] K. Simonyan and A. Zisserman, "Two-stream Convolutional Networks for Action Recognition in Videos," Neural Information Processing Systems, 2014.
- [195] S. Feichtinger, H. Weickert, and J. Weickert, "Edge-Enhancing Diffusion Filtering," IEEE Transactions on Pattern Analysis and Machine Intelligence, 2001.
- [196] C. Dwork and A. Roth, "The Algorithmic Foundations of Differential Privacy," Foundations and Trends in Theoretical Computer Science, 2014.
- [197] M. Hardt, K. Ligett, and F. McSherry, "The Reusable Holdout: Preserving Validity in Adaptive Data Analysis," Communications of the ACM, 2016.
- [198] J. Karla and A. Reiter, "Using Differential Privacy to Preserve Accuracy while Preventing Re-identification," Annual Privacy Law Workshop, 2017.
- [199] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative Adversarial Nets," Neural Information Processing Systems, 2014.
- [200] T. Karras, S. Laine, and T. Aila, "A Style-Based Generator Architecture for Generative Adversarial Networks," IEEE International Conference on Computer Vision, 2019.
- [201] X. Glorot and Y. Bengio, "Understanding the difficulty of training deep feedforward neural networks," International Conference on Artificial Intelligence and Statistics, 2010.
- [202] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," IEEE Conference on Computer Vision and Pattern Recognition, 2016.
- [203] B. Zhou, A. Lapedriza, J. Xiao, A. Torralba, and A. Oliva, "Learning Deep Features for Discriminative Localization," IEEE Conference on Computer Vision and Pattern Recognition, 2016.
- [204] S. Bae and P. Sonntag, "Edge-Preserving Image Super-Resolution from Multiple Images," IEEE Conference on Computer Vision and Pattern Recognition, 2013.
- [205] E. Shmueli, Y. Vaya, and E. Toch, "Towards Measuring Privacy," arXiv preprint arXiv:1410.4923, 2014.
- [206] L. Sweeney, "Simple Demographics Often Identify People Uniquely," Carnegie Mellon University, 2000.

- [207] A. Borji, D. N. Sihite, and L. Itti, "What Stands Out in a Crowd? Synergistic Effects of Population and Stimulus Driven Attention," *Journal of Vision*, 2012.
- [208] M. Cornia, L. Baraldi, G. Serra, and R. Cucchiara, "Predicting Human Eye Fixations via Saliency Maps," *IEEE Transactions on Image Processing*, 2016.
- [209] P. Isola, J. Xiao, A. Torralba, and A. Oliva, "What Makes an Image Memorable?" *IEEE Conference on Computer Vision and Pattern Recognition*, 2011.
- [210] A. Narayanan and V. Shmatikov, "Robust De-anonymization of Large Sparse Datasets," *IEEE Symposium on Security and Privacy*, 2008.
- [211] C. Dwork, "Differential Privacy: A Survey of Results," *Theory and Applications of Models of Computation*, 2008.
- [212] C. Dwork and A. Roth, "The Algorithmic Foundations of Differential Privacy," *Foundations and Trends in Theoretical Computer Science*, 2014.
- [213] C. E. Shannon, "A Mathematical Theory of Communication," *Bell System Technical Journal*, 1948.
- [214] T. M. Cover and J. A. Thomas, "Elements of Information Theory," *John Wiley & Sons*, 2006.
- [215] H. Nissenbaum, "Privacy in Context," *Stanford University Press*, 2009.
- [216] H. Nissenbaum, "Values in Design: Theory and Practice," *Springer*, 2016.
- [217] D. Boyd, "Networked Privacy," *Wiley Interdisciplinary Reviews: Internet Science*, 2014.
- [218] A. Squicciarini, M. Shehab, and F. Paci, "Collective Privacy Management in Social Networks," *Workshop on Online Social Networks*, 2009.
- [219] B. Shen and X. Liang, "Building Trust in the Era of Personalization," *HCI International Conference*, 2013.
- [220] J. Huang et al., "Speed/accuracy trade-offs for modern convolutional object detectors," *IEEE Conference on Computer Vision and Pattern Recognition*, 2017.
- [221] S. Ren, K. He, R. Girshick, and J. Sun, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2017.

- [222] Z. Zou et al., "Object Detection in 20 Years: A Survey," *Proceedings of the IEEE*, 2023.
- [223] K. He, G. Gkioxari, P. Dollár, and R. Girshick, "Mask R-CNN," *IEEE International Conference on Computer Vision*, 2017.
- [224] R. Girshick, J. Donahue, T. Darrell, and J. Malik, "Rich Feature Hierarchies for Accurate Object Detection and Semantic Segmentation," *IEEE Conference on Computer Vision and Pattern Recognition*, 2014.
- [225] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet Classification with Deep Convolutional Neural Networks," *Communications of the ACM*, 2017.
- [226] K. Simonyan and A. Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition," *International Conference on Learning Representations*, 2015.
- [227] Y. Zhou, S.-I. Morikawa, and K. Yamanaka, "Quantized Neural Networks: Training Neural Networks with Low Precision Weights and Activations," *arXiv preprint arXiv:1609.07061*, 2016.
- [228] M. Courbariaux, I. Hubara, and Y. Bengio, "Binarized Neural Networks," *Neural Information Processing Systems*, 2016.
- [229] S. Han, H. Mao, and W. J. Dally, "Deep Compression: Compressing Deep Neural Networks with Pruning, Trained Quantization and Huffman Coding," *International Conference on Learning Representations*, 2016.
- [230] Y. LeCun, J. S. Denker, and S. A. Solla, "Optimal Brain Damage," *Neural Information Processing Systems*, 1990.
- [231] G. Hinton, O. Vinyals, and J. Dean, "Distilling the Knowledge in a Neural Network," *Neural Information Processing Systems Workshop*, 2015.
- [232] J. Yim, D. Joo, B. Bae, and J. Kim, "A Gift from Knowledge Distillation: Fast Optimization, Network Minimization and Transfer Learning," *IEEE Conference on Computer Vision and Pattern Recognition*, 2017.
- [233] S. Mohseni, N. Zarei, and E. D. Ragan, "A Multidisciplinary Survey and Framework for Design and Evaluation of Explainable AI Systems," *ACM Transactions on Interactive Intelligent Systems*, 2021.

- [234] B. Kim, M. Wattenberg, J. Gilmer, C. Cai, J. Wexler, F. Viégas, and R. Sayres, "Interpretability Beyond Feature Attribution: Quantitative Testing with Concept Activation Vectors," International Conference on Machine Learning, 2018.
- [235] F. Schaub, R. Balebako, A. L. Durity, and L. F. Cranor, "A Design Space for Effective Privacy Notices," Symposium on Usable Privacy and Security, 2015.
- [236] L. A. Zuboff, "Surveillance Capitalism: The Fight for a Human Future at the New Frontier of Power," PublicAffairs, 2019.
- [237] Y. Wang et al., "The Effects of Individual Differences on Security Notifications," Symposium on Usable Privacy and Security, 2014.
- [238] J. Staddon et al., "A Review of the Usable Privacy and Security Literature," Technical Report CMU-ISR-10-122, Carnegie Mellon University, 2012.
- [239] N. Patrignani and D. Whitehouse, "Slow Tech: Critical Reflection on Technology and Slow Living," Palgrave Macmillan, 2018.
- [240] H. T. Tavani, "Understanding Computing Ethics: Privacy, Information Ethics and Policy," arXiv preprint, 2006.
- [241] J. Deng et al., "RetinaFace: Single-stage Dense Face Localisation in the Wild," IEEE Conference on Computer Vision and Pattern Recognition, 2020.
- [242] R. Laroca et al., "A Robust Real-Time Automatic License Plate Recognition Based on Deep Convolutional Neural Networks," IEEE Transactions on Intelligent Transportation Systems, 2021.
- [243] O. Russakovsky, J. Deng, H. Su, J. Krause, S. Satheesh, S. Ma, Z. Huang, A. Karpukhin, A. Khosla, M. Bernstein, A.-L. Fei-Fei, and L. Fei-Fei, "ImageNet Large Scale Visual Recognition Challenge," International Journal of Computer Vision, 2015.
- [244] B. Shi, M. Yang, X. Wang, P. Lyu, and C. Yao, "CRNN: An End-to-End Trainable Neural Network for Scene Text Recognition," International Conference on Computer Vision, 2017.
- [245] Z. Wang, A. C. Bovik, H. R. Sheikh, and E. P. Simoncelli, "Image Quality Assessment: From Error Visibility to Structural Similarity," IEEE Transactions on Image Processing, 2004.

- [246] M. Courbariaux, Y. Bengio, and A. Courville, "Binarized Convolutional Neural Networks with Separable Filters for Efficient Hardware Acceleration," IEEE International Conference on Computer Vision Workshops, 2015.
- [247] T. Karras, S. Laine, and T. Aila, "A Style-Based Generator Architecture for Generative Adversarial Networks," IEEE International Conference on Computer Vision, 2019.
- [248] J. Huang et al., "Speed/accuracy trade-offs for modern convolutional object detectors," IEEE Conference on Computer Vision and Pattern Recognition, 2017.
- [249] G. Jocher, A. Chaurasia, and A. Qiu, "YOLO by Ultralytics," GitHub Repository, 2023.
- [250] N. Ma, X. Zhang, H. Zheng, and J. Sun, "ShuffleNet V2: Practical Guidelines for Efficient CNN Architecture Design," European Conference on Computer Vision, 2018.
- [251] T. Xu, P. Zhang, Q. Huang, H. Zhang, Z. Gan, X. Huang, and X. He, "AttnGAN: Fine-Grained Text to Image Generation with Attentional Generative Network," IEEE International Conference on Computer Vision, 2018.
- [252] Y. Wang et al., "The Effects of Individual Differences on Security Notifications," Symposium on Usable Privacy and Security, 2014.
- [253] F. Schaub, R. Balebako, A. L. Durity, and L. F. Cranor, "A Design Space for Effective Privacy Notices," Symposium on Usable Privacy and Security, 2015.
- [254] B. Kim, M. Wattenberg, J. Gilmer, C. Cai, J. Wexler, F. Viégas, and R. Sayres, "Interpretability Beyond Feature Attribution: Quantitative Testing with Concept Activation Vectors," International Conference on Machine Learning, 2018.
- [255] V. Roth and M. Greiff, "The Other Side of Accuracy: Privacy Loss in Machine Learning," arXiv preprint arXiv:2006.09109, 2020.
- [256] C. Dwork and A. Roth, "The Algorithmic Foundations of Differential Privacy," Foundations and Trends in Theoretical Computer Science, 2014.
- [257] H. P. Kramer, "Low-light Image Enhancement via a Deep Hybrid Network," IEEE Transactions on Image Processing, 2019.
- [258] Y. Ding, X. Zhang, L. Zhang, and D. Doermann, "Scene Text Detection and Recognition: Recent Advances and Future Trends," International Journal on Document

Analysis and Recognition, 2015.

[259] B. D. Lucas and T. Kanade, "An Iterative Image Registration Technique with an Application to Stereo Vision," International Joint Conference on Artificial Intelligence, 1981.

[260] M. Patil, A. Westin, G. Wills, and A. Forsgren, "Privacy-Preserving Mechanisms for Supporting Dynamic Collaborative Multi-domain Environments," Privacy Technology and Policy Workshop, 2010.

[261] S. Mohseni, N. Zarei, and E. D. Ragan, "A Multidisciplinary Survey and Framework for Design and Evaluation of Explainable AI Systems," ACM Transactions on Interactive Intelligent Systems, 2021.

[262] C. Szegedy, W. Zaremba, I. Sutskever, J. Bruna, D. Erhan, I. Goodfellow, and R. Fergus, "Intriguing properties of neural networks," International Conference on Learning Representations, 2014.

[263] I. J. Goodfellow, J. Shlens, and C. Szegedy, "Explaining and Harnessing Adversarial Examples," International Conference on Learning Representations, 2015.

[264] M. Carlini and D. Wagner, "Towards Evaluating the Robustness of Neural Networks," IEEE Symposium on Security and Privacy, 2017.

[265] A. Narayanan and V. Shmatikov, "Robust De-anonymization of Large Sparse Datasets," IEEE Symposium on Security and Privacy, 2008.

[266] L. Sweeney, "Simple Demographics Often Identify People Uniquely," Carnegie Mellon University, 2000.

[267] T. Lipton and J. Steinhardt, "Counterexample Guided Abstraction Refinement," arXiv preprint arXiv:1910.13699, 2019.

[268] S. Ren, K. He, R. Girshick, and J. Sun, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks," IEEE Transactions on Pattern Analysis and Machine Intelligence, 2017.

[269] A. Dosovitskiy et al., "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale," International Conference on Learning Representations, 2021.

[270] N. Carion, F. Massa, G. Synnaeve, N. Usunier, A. Kirillov, and S. Zagoruyko, "End-to-End Object Detection with Transformers," European Conference on Computer

Vision, 2020.

[271] A. Vaswani et al., "Attention is All You Need," Neural Information Processing Systems, 2017.

[272] Z. Liu, Y. Lin, Y. Cao, H. Hu, Y. Wei, Z. Zhang, S. Lin, and B. Guo, "Swin Transformer: Hierarchical Vision Transformer using Shifted Windows," IEEE International Conference on Computer Vision, 2021.

[273] X. Wang, R. Girshick, A. Gupta, and K. He, "Non-local Neural Networks," IEEE Conference on Computer Vision and Pattern Recognition, 2018.

[274] J.-B. Alayrac, B. Recasens, R. Schneider, R. Arandjelović, J. Ramapuram, J. De Fauw, L. Smaira, S. Dieleman, and A. Zisserman, "Self-Supervised MultiModal Learning from Video," arXiv preprint arXiv:2104.14407, 2021.

[275] B. Z. Yuksekogonul et al., "When and Why is a Concept Activation Vector Useful? Quantifying Post-Hoc Interpretability," Neural Information Processing Systems, 2022.

[276] T. Brown et al., "Language Models are Few-Shot Learners," Neural Information Processing Systems, 2020.

[277] C. Finn, P. Abbeel, and S. Levine, "Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks," International Conference on Machine Learning, 2017.

[278] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-Efficient Learning of Deep Networks from Decentralized Data," Artificial Intelligence and Statistics, 2017.

[279] S. Caldas et al., "Leaf: A Benchmark for Federated Settings," arXiv preprint arXiv:1812.01097, 2018.

[280] T. Li, A. K. Sahu, M. Zaheer, M. Srivastava, A. Talwalkar, and V. Smith, "Federated Optimization in Heterogeneous Networks," Machine Learning Systems Workshop, 2018.

[281] C. Dwork, "Differential Privacy: A Survey of Results," Theory and Applications of Models of Computation, 2008.

[282] M. Hardt, K. Ligett, and F. McSherry, "The Reusable Holdout: Preserving Validity in Adaptive Data Analysis," Communications of the ACM, 2016.

- [283] J. Karla and A. Reiter, "Using Differential Privacy to Preserve Accuracy while Preventing Re-identification," Annual Privacy Law Workshop, 2017.
- [284] R. L. Rivest, L. Adleman, and M. L. Dertouzos, "On Data Banks and Privacy Homomorphisms," Foundations of Secure Computation, 1978.
- [285] Z. Brakerski, C. Gentry, and V. Vaikuntanathan, "Fully Homomorphic Encryption without Bootstrapping," International Conference on Cryptography, 2012.
- [286] D. Archer, S. Brave, W. Butcher, J. Catalano, B. Hales, F. Kerschbaum, W. Künzle, K. Leontjev, R. Ostrovsky, and D. Schuster, "Recommendations for Securing Fully Homomorphic Encryption," arXiv preprint arXiv:1909.01108, 2019.
- [287] B. Kim, M. Wattenberg, J. Gilmer, C. Cai, J. Wexler, F. Viégas, and R. Sayres, "Interpretability Beyond Feature Attribution: Quantitative Testing with Concept Activation Vectors," International Conference on Machine Learning, 2018.
- [288] M. T. Ribeiro, S. Singh, and C. Guestrin, "Why should I trust you?: Explaining the predictions of any classifier," International Conference on Knowledge Discovery and Data Mining, 2016.
- [289] A. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, "Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization," IEEE International Conference on Computer Vision, 2017.
- [290] J. Simonyan, A. Vedaldi, and A. Zisserman, "Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps," ICLR Workshop, 2014.
- [291] B. Lepri, N. Oliver, E. Letouzé, A. Pentland, and P. Vinck, "Fair, Transparent, and Accountable Algorithmic Decision-making Processes," Philosophy & Technology, 2018.
- [292] S. Wachter, B. Mittelstadt, and L. Floridi, "Why a Right to Explanation of Automated Decision-making Does Not Exist in the General Data Protection Regulation," International Data Privacy Law, 2017.
- [293] B. Barocas and K. Selbst, "Big Data's Disparate Impact," California Law Review, 2016.
- [294] S. Buolamwini and T. Gebru, "Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification," Conference on Fairness, Accountability and Transparency, 2018.

- [295] M. Hardt, E. Price, and N. Srebro, "Equality of Opportunity in Supervised Learning," Neural Information Processing Systems, 2016.
- [296] A. Morales-Guzman, H. Nissenbaum, and D. Hendler, "Contextual Integrity for Privacy and Data Protection," arXiv preprint, 2021.
- [297] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative Adversarial Nets," Neural Information Processing Systems, 2014.
- [298] T. Karras, S. Laine, and T. Aila, "A Style-Based Generator Architecture for Generative Adversarial Networks," IEEE International Conference on Computer Vision, 2019.
- [299] A. Radford, L. Metz, and S. Chintala, "Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks," International Conference on Learning Representations, 2016.
- [300] J. Ho, A. Jain, and P. Abbeel, "Denoising Diffusion Probabilistic Models," Neural Information Processing Systems, 2020.
- [301] J. Song, C. Meng, and S. Ermon, "Denoising Diffusion Implicit Models," International Conference on Learning Representations, 2021.
- [302] Y. Song and S. Ermon, "Generative Modeling by Estimating Gradients of the Data Distribution," Neural Information Processing Systems, 2019.
- [303] D. P. Kingma and M. Welling, "Auto-Encoding Variational Bayes," International Conference on Learning Representations, 2014.
- [304] I. Higgins, L. Matthey, A. Pal, C. Burgess, X. Glorot, M. Botvinick, S. Mohamed, and Y. Lerchner, " β -VAE: Learning Basic Visual Concepts with a Constrained Variational Framework," International Conference on Learning Representations, 2017.
- [305] European Union, "Regulation (EU) 2016/679 of the European Parliament and of the Council," Official Journal of the European Union, 2016.
- [306] V. Leavitt, "Meta Faces Record Privacy Fine for Cookie Violations," Law360, 2022.
- [307] L. A. Bygrave, "Data Protection by Design and by Default: Deciphering the EU's Legislative Endeavours," European Data Protection Law Review, 2014.

- [308] H. Kraemer, D. C. Coleman, and J. E. Stahl, "Privacy by Design," arXiv preprint, 2015.
- [309] T. Cavoukian, "Privacy by Design: The 7 Foundational Principles," Information & Privacy Commissioner, 2011.
- [310] B. Mahoney, "The Evolution of Data Protection Law in the Cloud," Journal of Internet Law, 2018.
- [311] P. Schwartz and K. Peifer, "Transatlantic Data Privacy," Georgetown Law Technology Review, 2017.
- [312] H. Nissenbaum, "Privacy in Context," Stanford University Press, 2009.
- [313] H. Nissenbaum, "Values in Design: Theory and Practice," Springer, 2016.
- [314] D. Boyd, "Networked Privacy," Wiley Interdisciplinary Reviews: Internet Science, 2014.
- [315] P. Schwartz and B. Solove, "PII, Privacy and Regulating Public Data," Georgetown Privacy Law, 2014.
- [316] N. Patil and C. Tuli, "Privacy in the Era of Big Data: A Systematic Overview," arXiv preprint, 2016.
- [317] A. Acquisti, L. F. Cranor, and R. E. Smith, "A Framework for Human Factors in Information Security," Workshop on New Security Paradigms, 2006.
- [318] F. Schaub, R. Balebako, A. L. Durity, and L. F. Cranor, "A Design Space for Effective Privacy Notices," Symposium on Usable Privacy and Security, 2015.
- [319] R. Gross, E. Acquisti, and H. F. Heinz, "Information revelation and privacy in online social networks," Workshop on Privacy in Electronic Society, 2005.
- [320] J. Staddon et al., "A Review of the Usable Privacy and Security Literature," Technical Report CMU-ISR-10-122, Carnegie Mellon University, 2012.
- [321] Z. Zou et al., "Object Detection in 20 Years: A Survey," Proceedings of the IEEE, 2023.
- [322] Y. Wang et al., "The Effects of Individual Differences on Security Notifications," Symposium on Usable Privacy and Security, 2014.

[323] V. Roth and M. Greiff, "The Other Side of Accuracy: Privacy Loss in Machine Learning," arXiv preprint arXiv:2006.09109, 2020.

[324] M. Al Faruque and S. Sibin, "Towards Secure Autonomous Cyber-Physical Interaction in Next Generation Vehicles," IEEE Asia Pacific Cloud Computing Conference, 2016.

[325] L. A. Zuboff, "Surveillance Capitalism: The Fight for a Human Future at the New Frontier of Power," PublicAffairs, 2019.

[326] D. Solove, "Understanding Privacy," Harvard University Press, 2008.

[327] A. Acquisti and J. Grossklags, "Uncertainty, Ambiguity and Privacy," Privacy Technology and Policy Workshop, 2007.

[328] European Union, "Regulation (EU) 2016/679 of the European Parliament and of the Council," Official Journal of the European Union, 2016.

[329] H. Nissenbaum, "Privacy in Context," Stanford University Press, 2009.

Document Status: Complete Survey Paper - Literature Review

Version: 1.0

Date: April 28, 2026